

Adventures in Concurrent Garbage Collection

**Partially based on slides or graphs by
Erik Österlund, Albert Mingkun Yang, Jonas Norlinder**

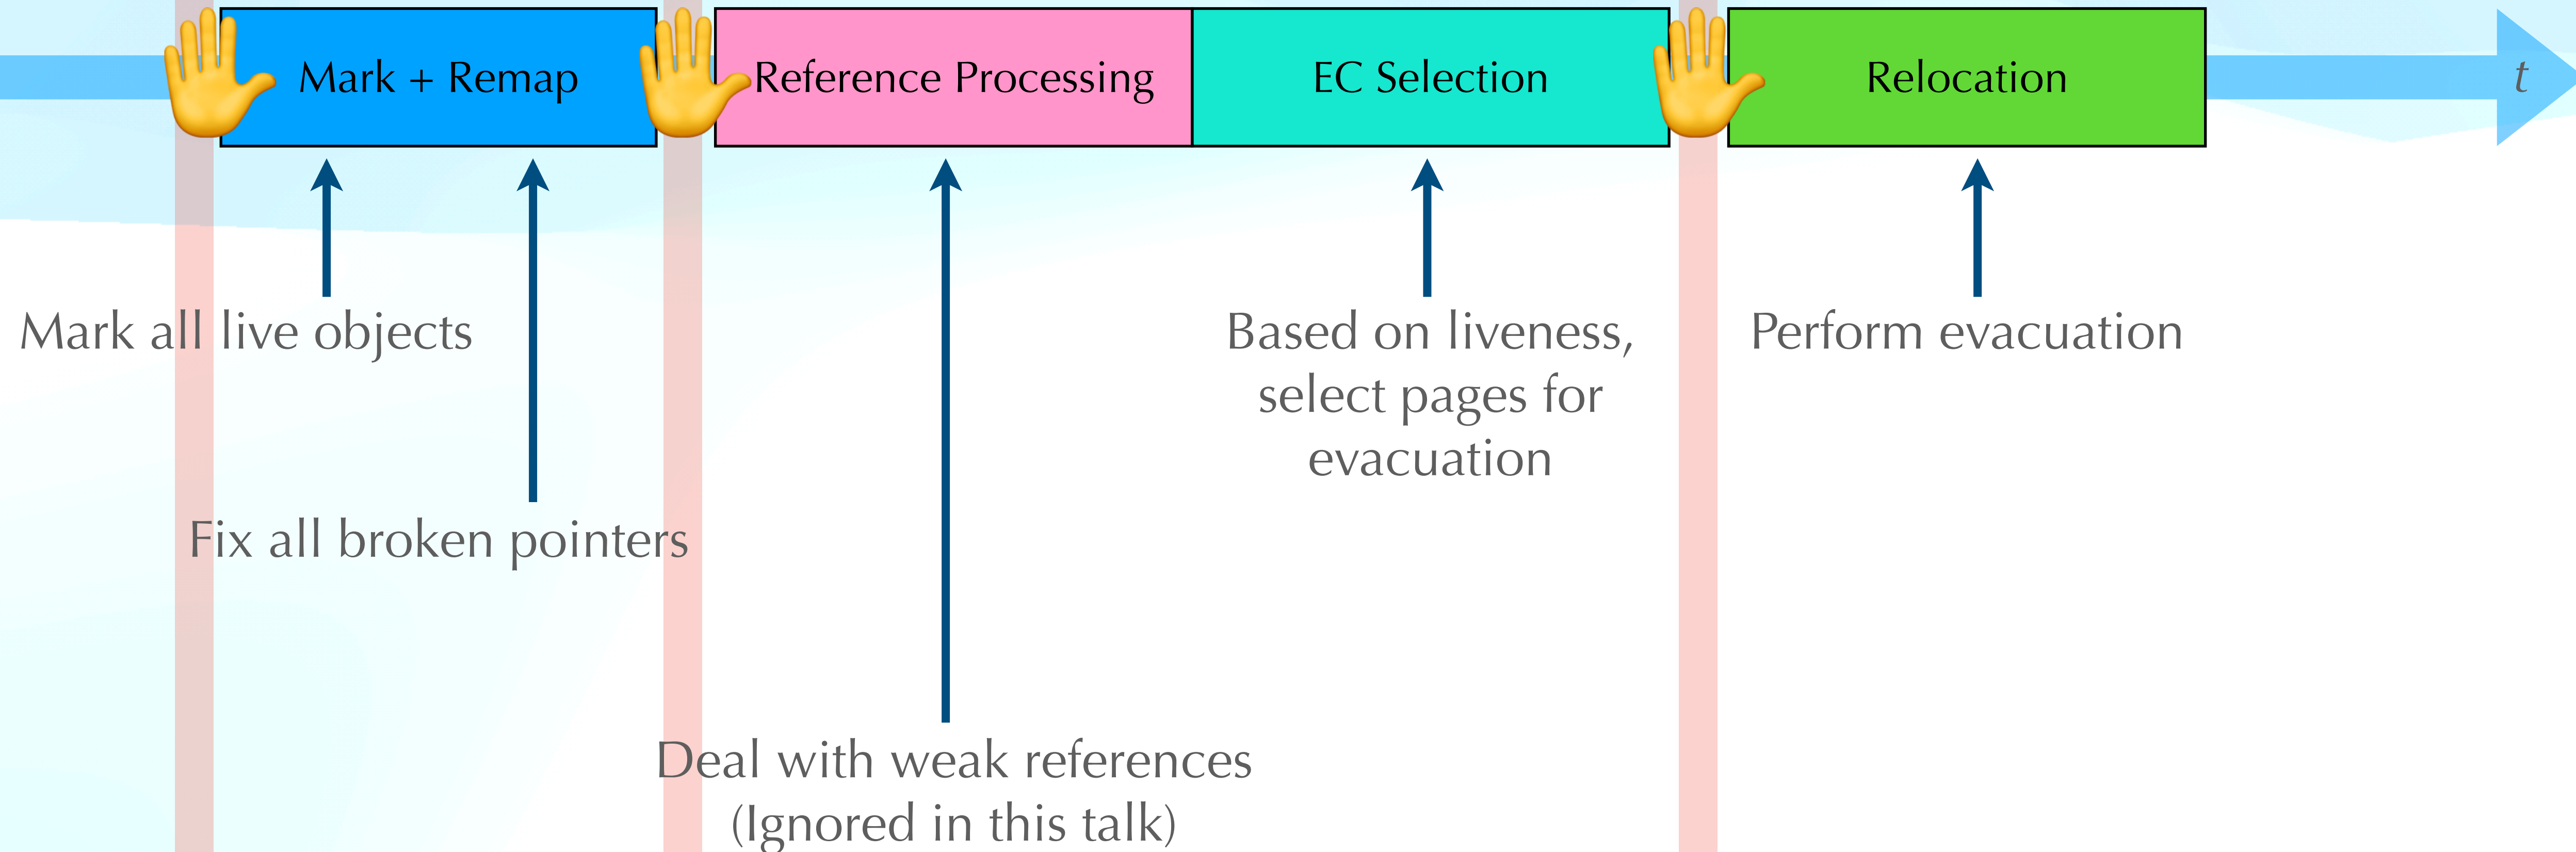
Tobias Wrigstad at PLISS/FOJW 2023

The Concurrent Z Garbage Collector (ZGC)

In OpenJDK since 11 (experimental); 15 (prod.); 21 (generational)

- **Goals:** tail-latency, TB size heaps, <1ms pause times, no need for tuning
- Allows GC to run concurrent with mutators
 - Global STW pauses whose lengths are invariant of heap size
- Concurrent marking (as explained to us by Tony yesterday)
- Concurrent compaction (as to be explained to us by Tony (and myself))
- Concurrent weak reference processing (ask Richard during coffee)
- Developed by Oracle (Per Lidén, Stefan Karlsson, Erik Österlund, et al.)

Overview of Single-Generation ZGC



Quick Load-Barrier Primer particular to ZGC

- ZGC is a concurrent GC in OpenJDK
 - GC threads are free to move objects around
 - **Program threads discover this in load barriers**
 - Load barrier slow paths heal dangling pointers
 - Load-barrier overhead ~2–3%

```
Object o = x.f
```

Load barrier needed


```
Object o = x
```

```
x.foo(y)
```

```
int i = x.f
```

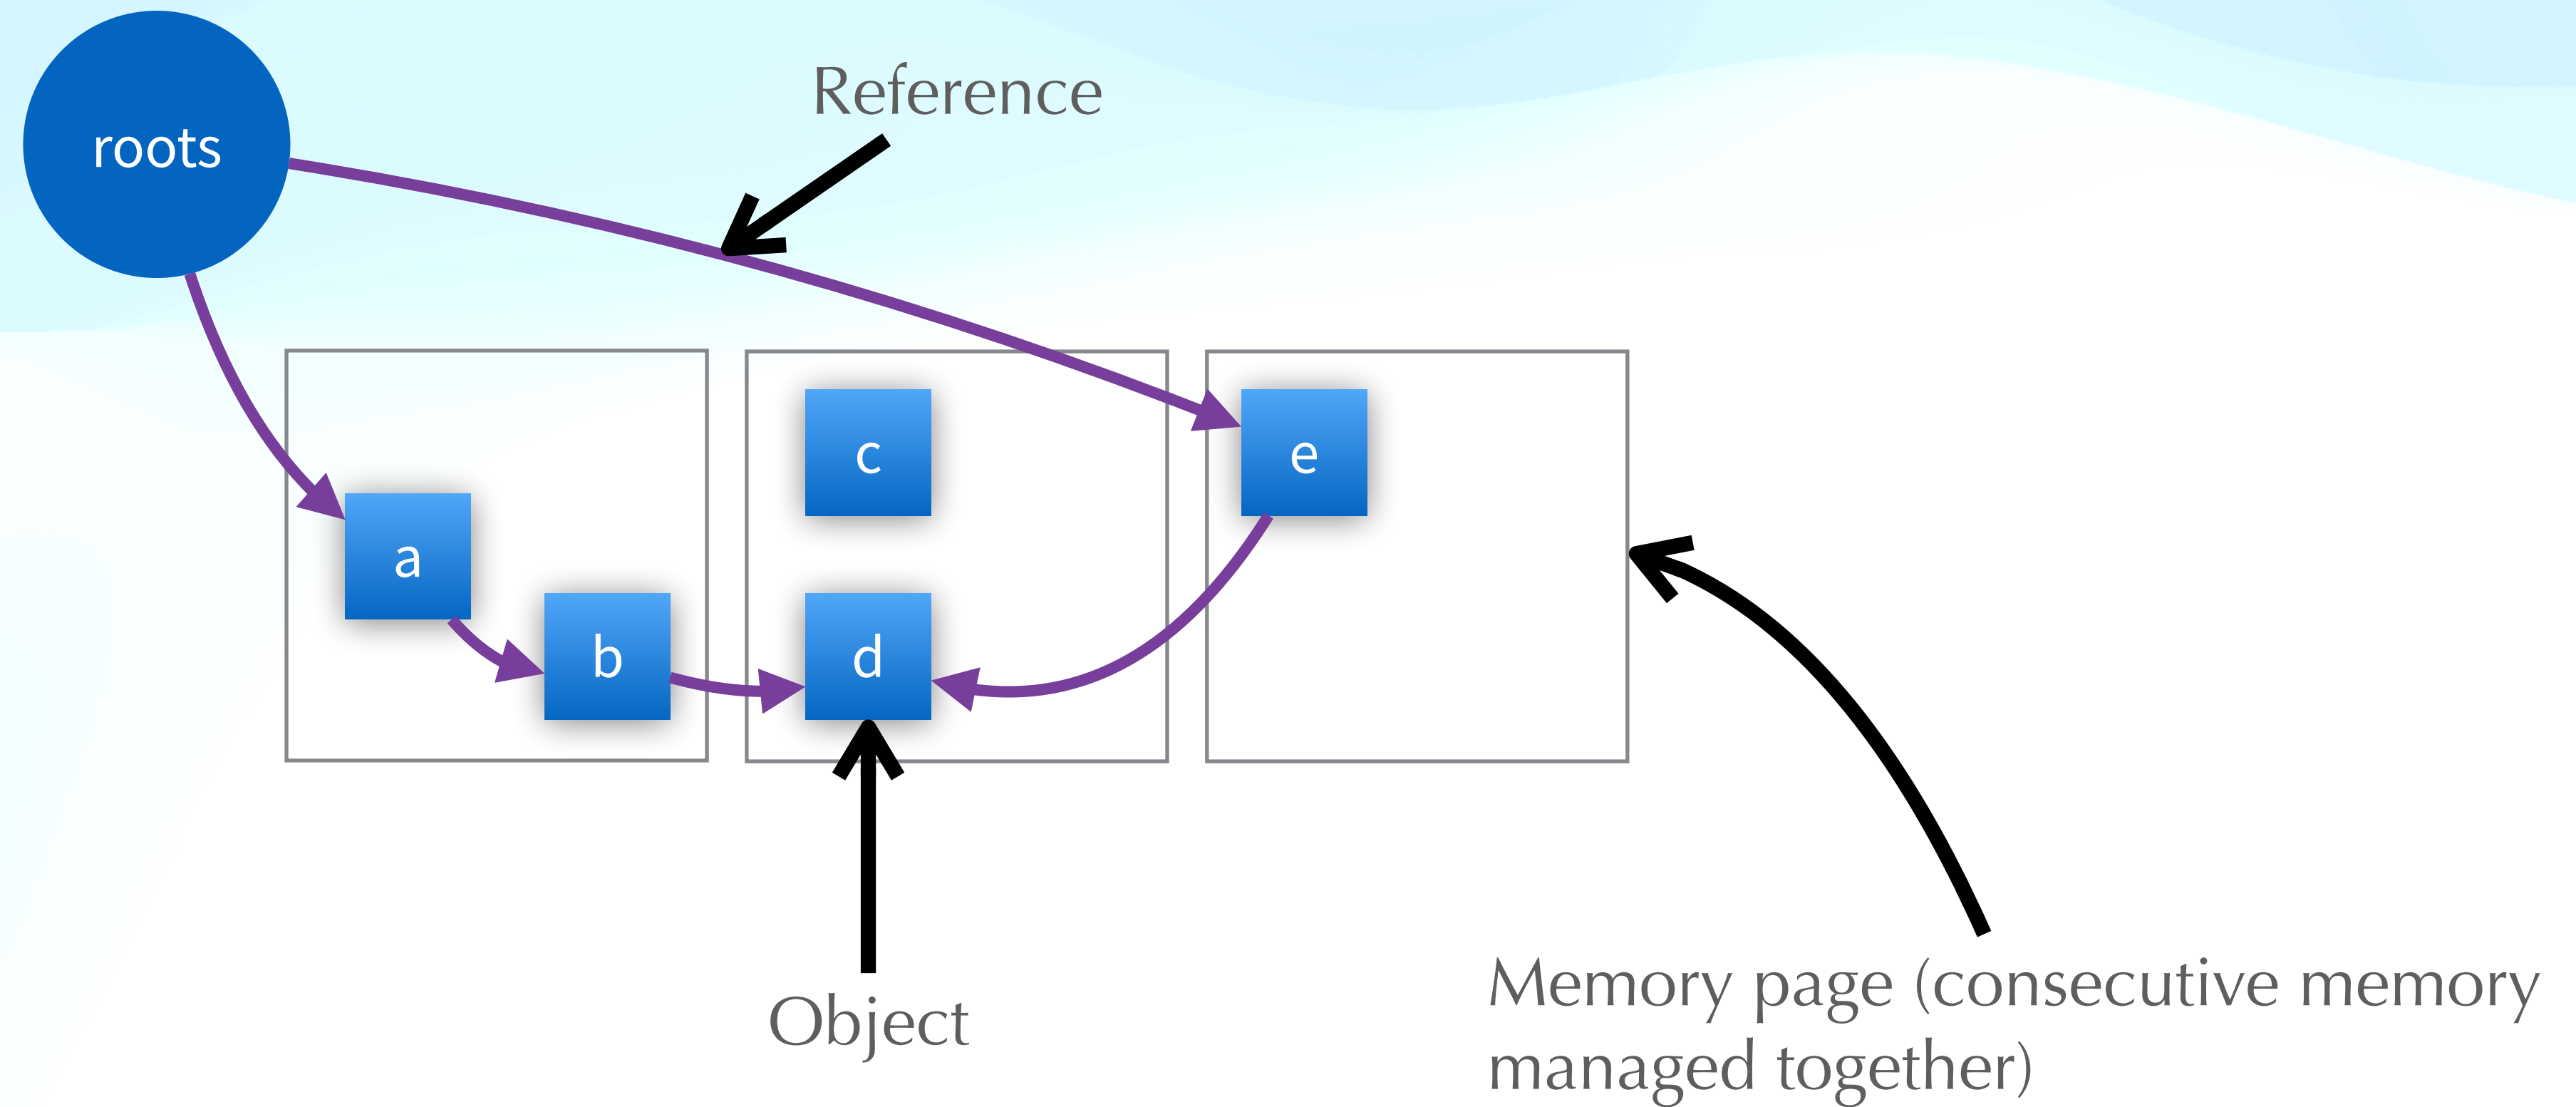
Load barrier not needed

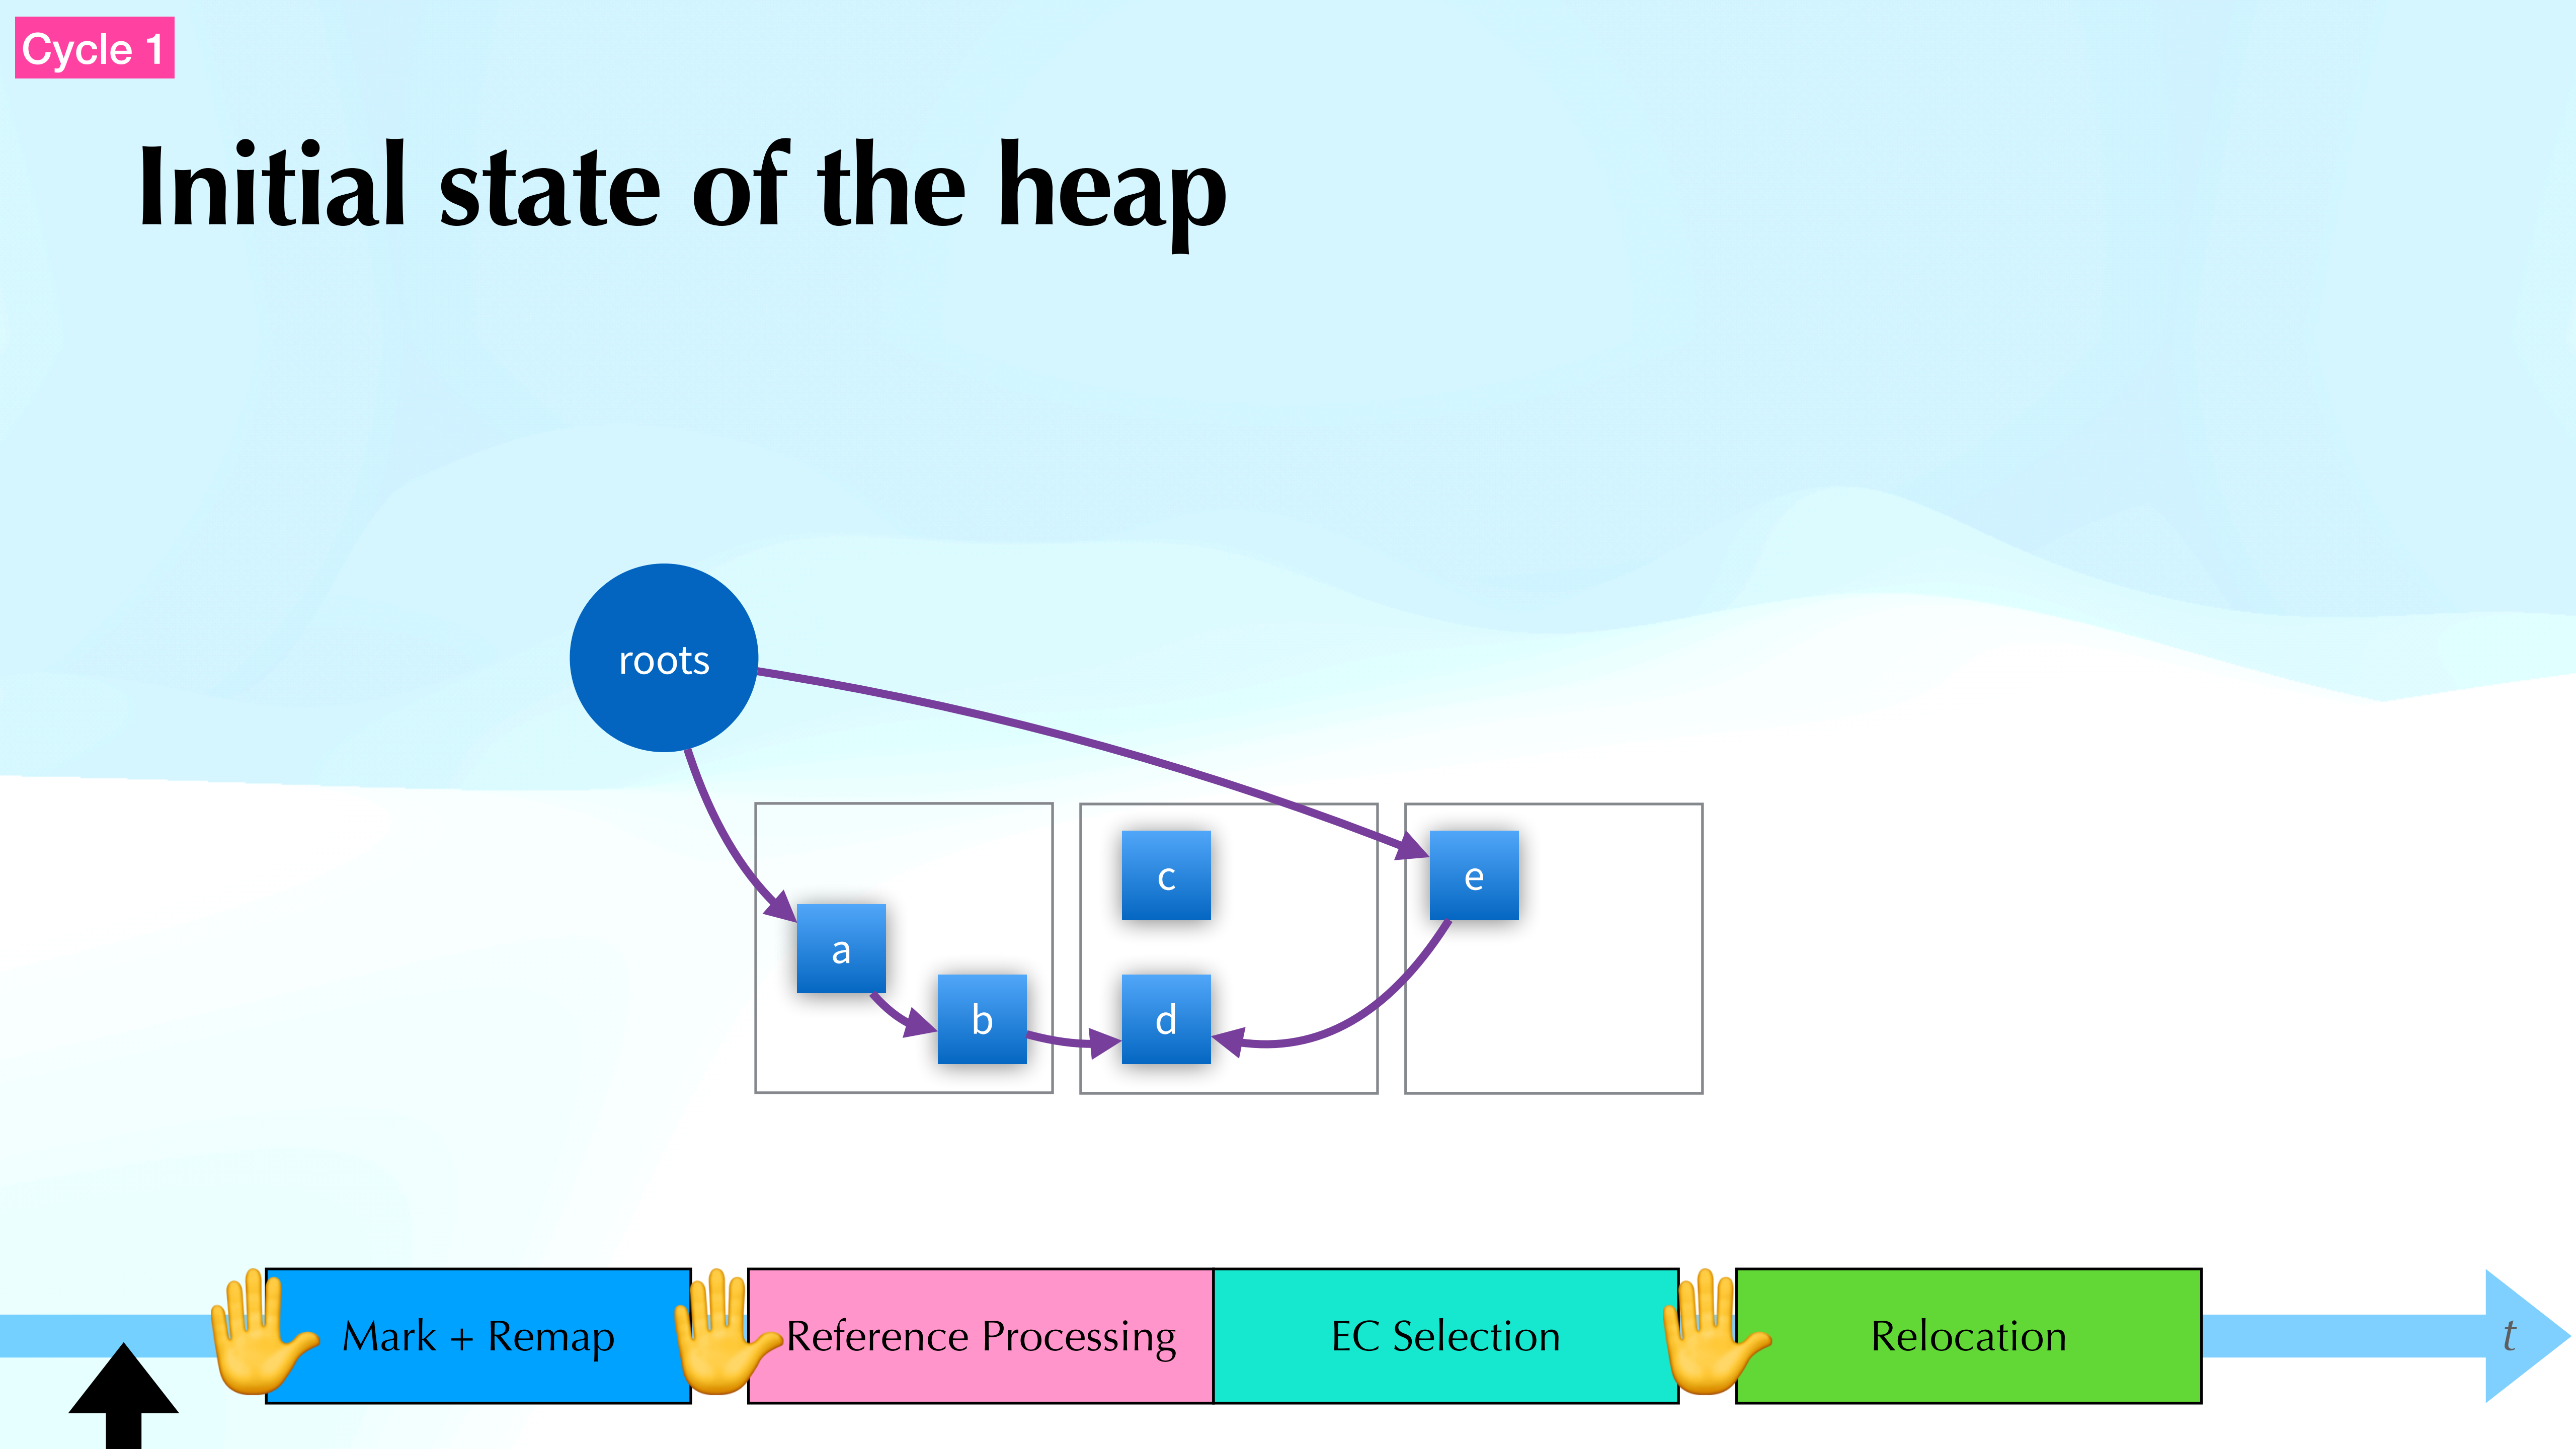
Anatomy of a pointer in ZGC

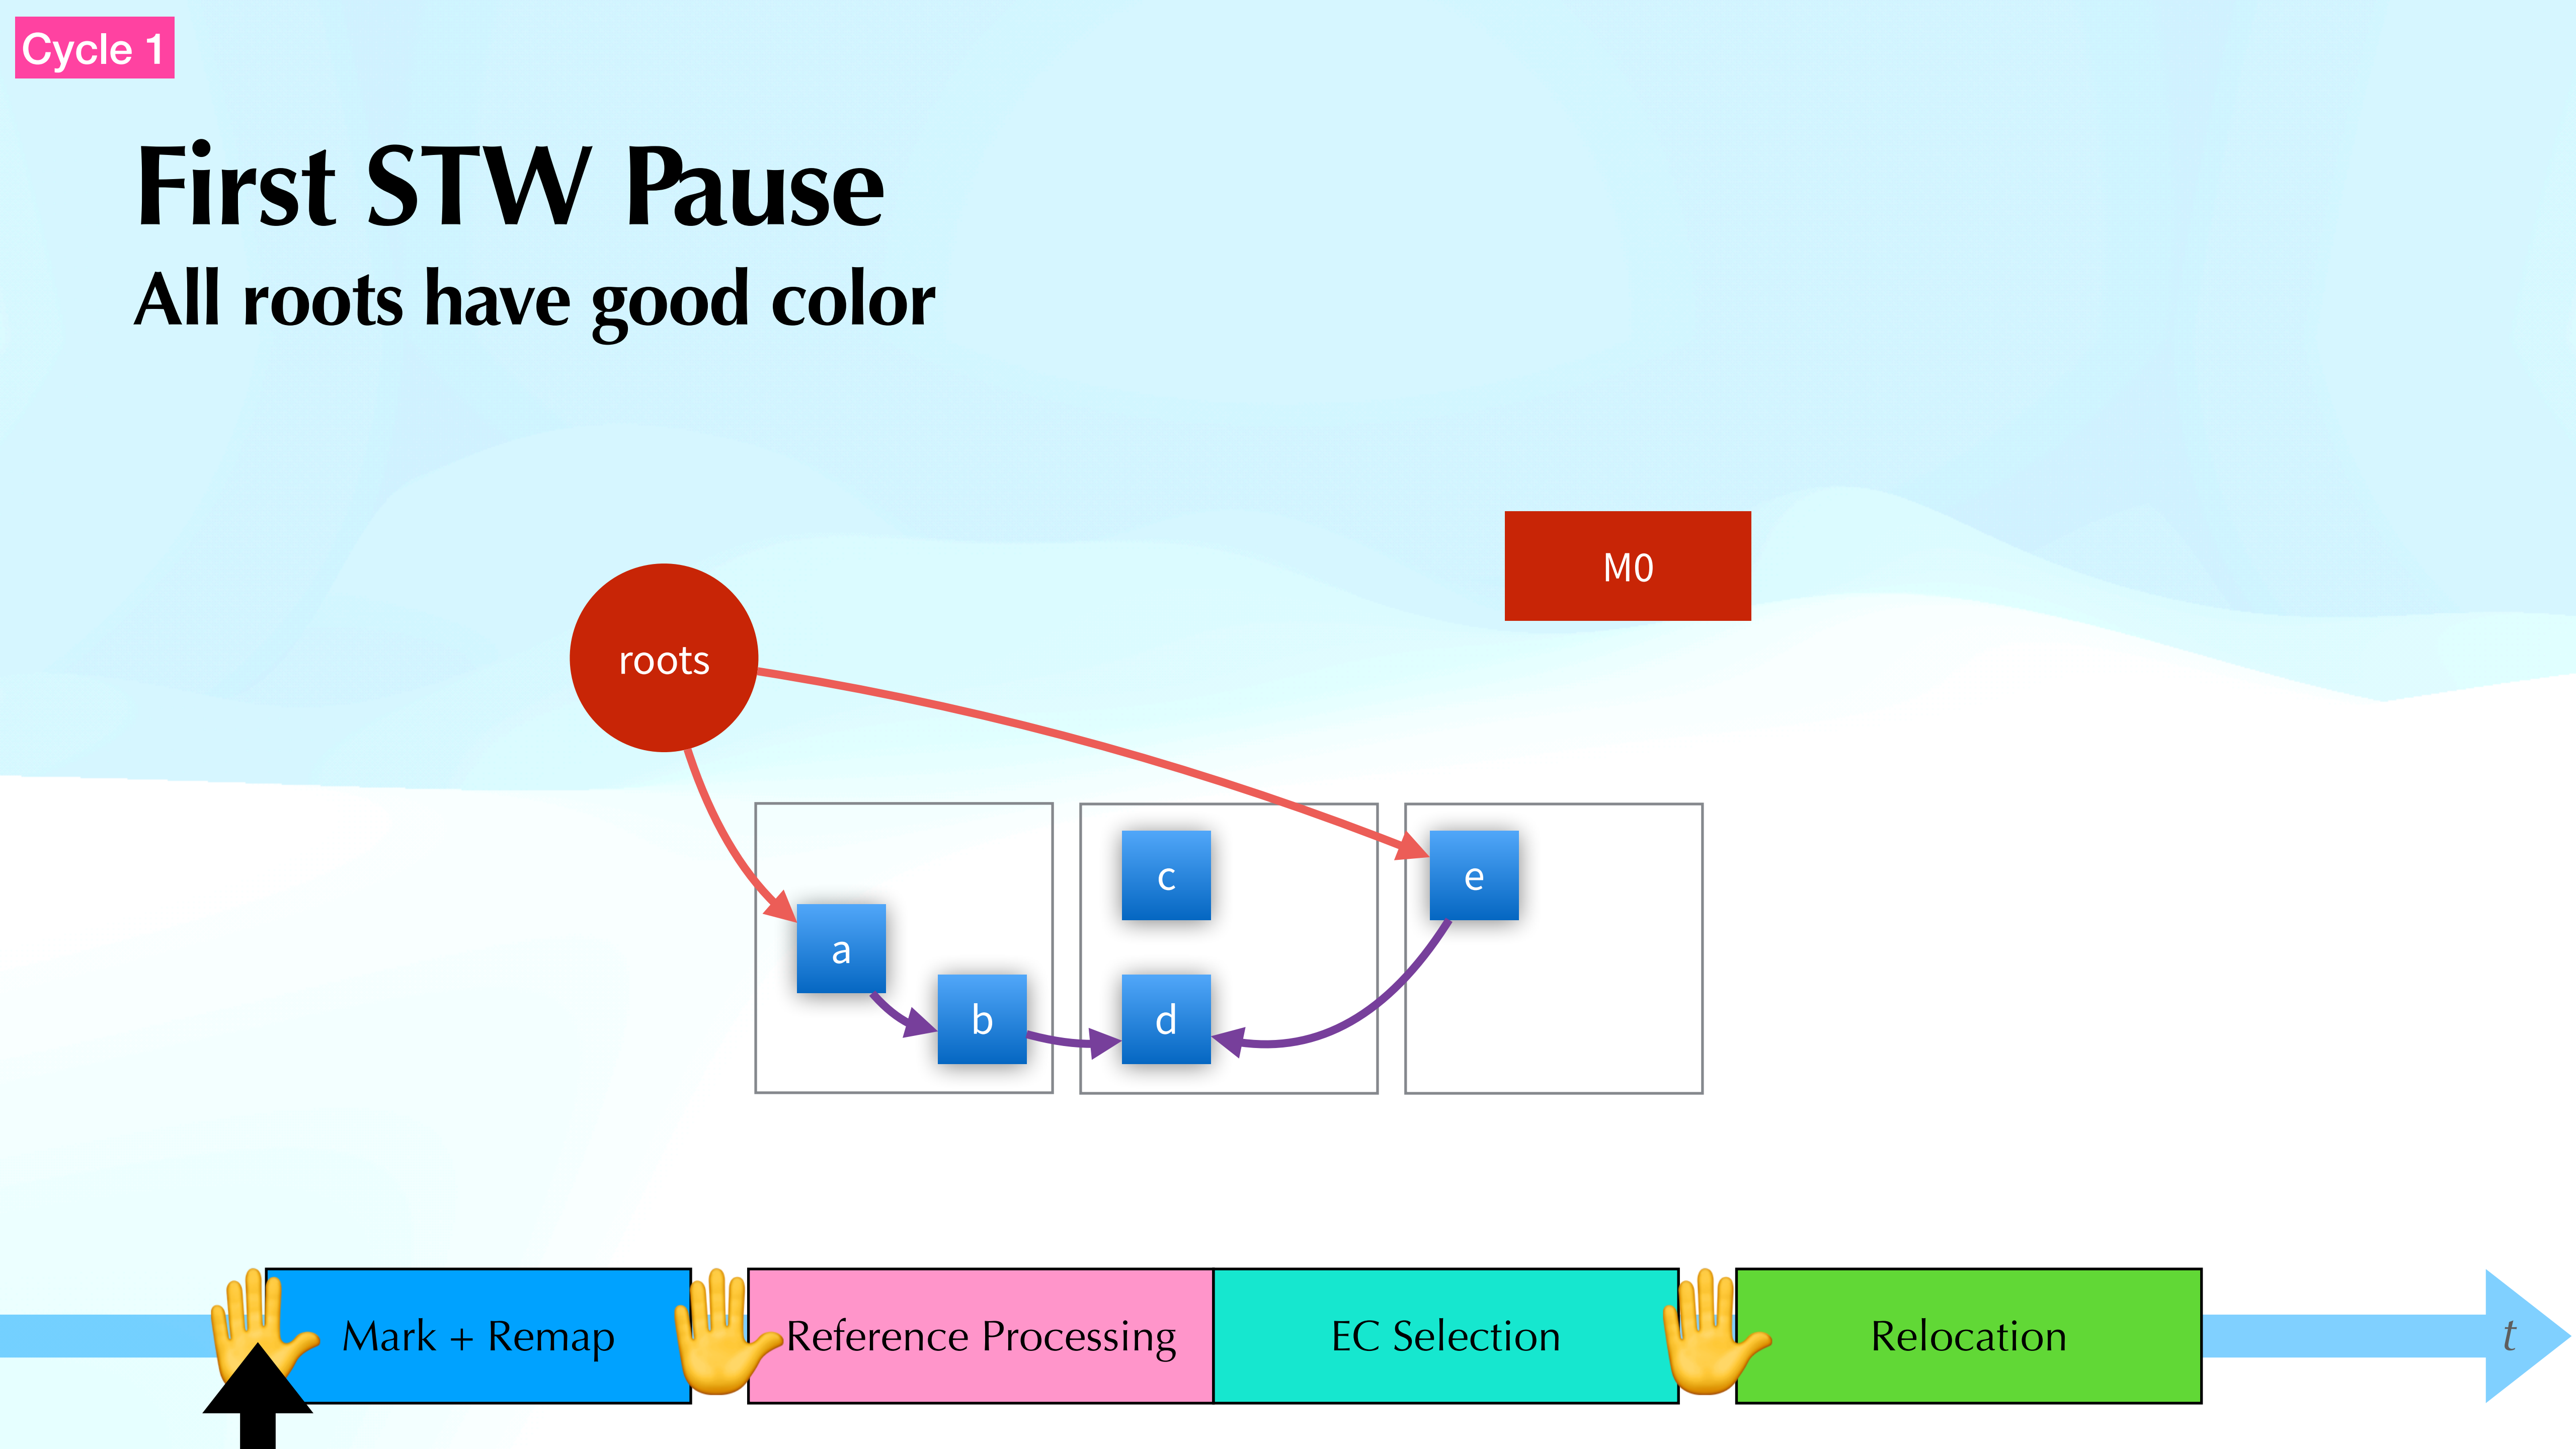


- Use a load barrier when storing addresses from the heap to the stack
- Good colour — fast path, do nothing
- Bad colour — slow path, do something

Example of ZGC Cycle

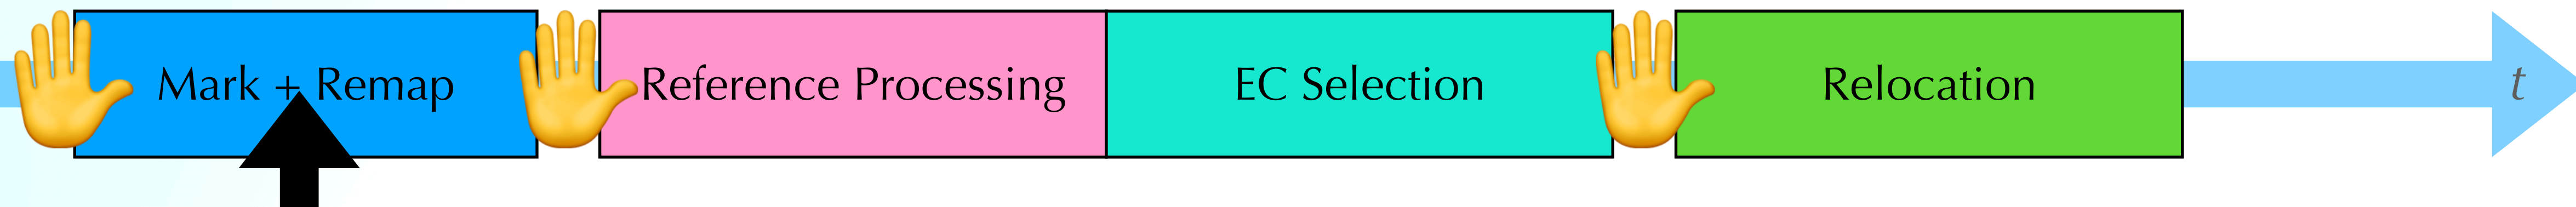
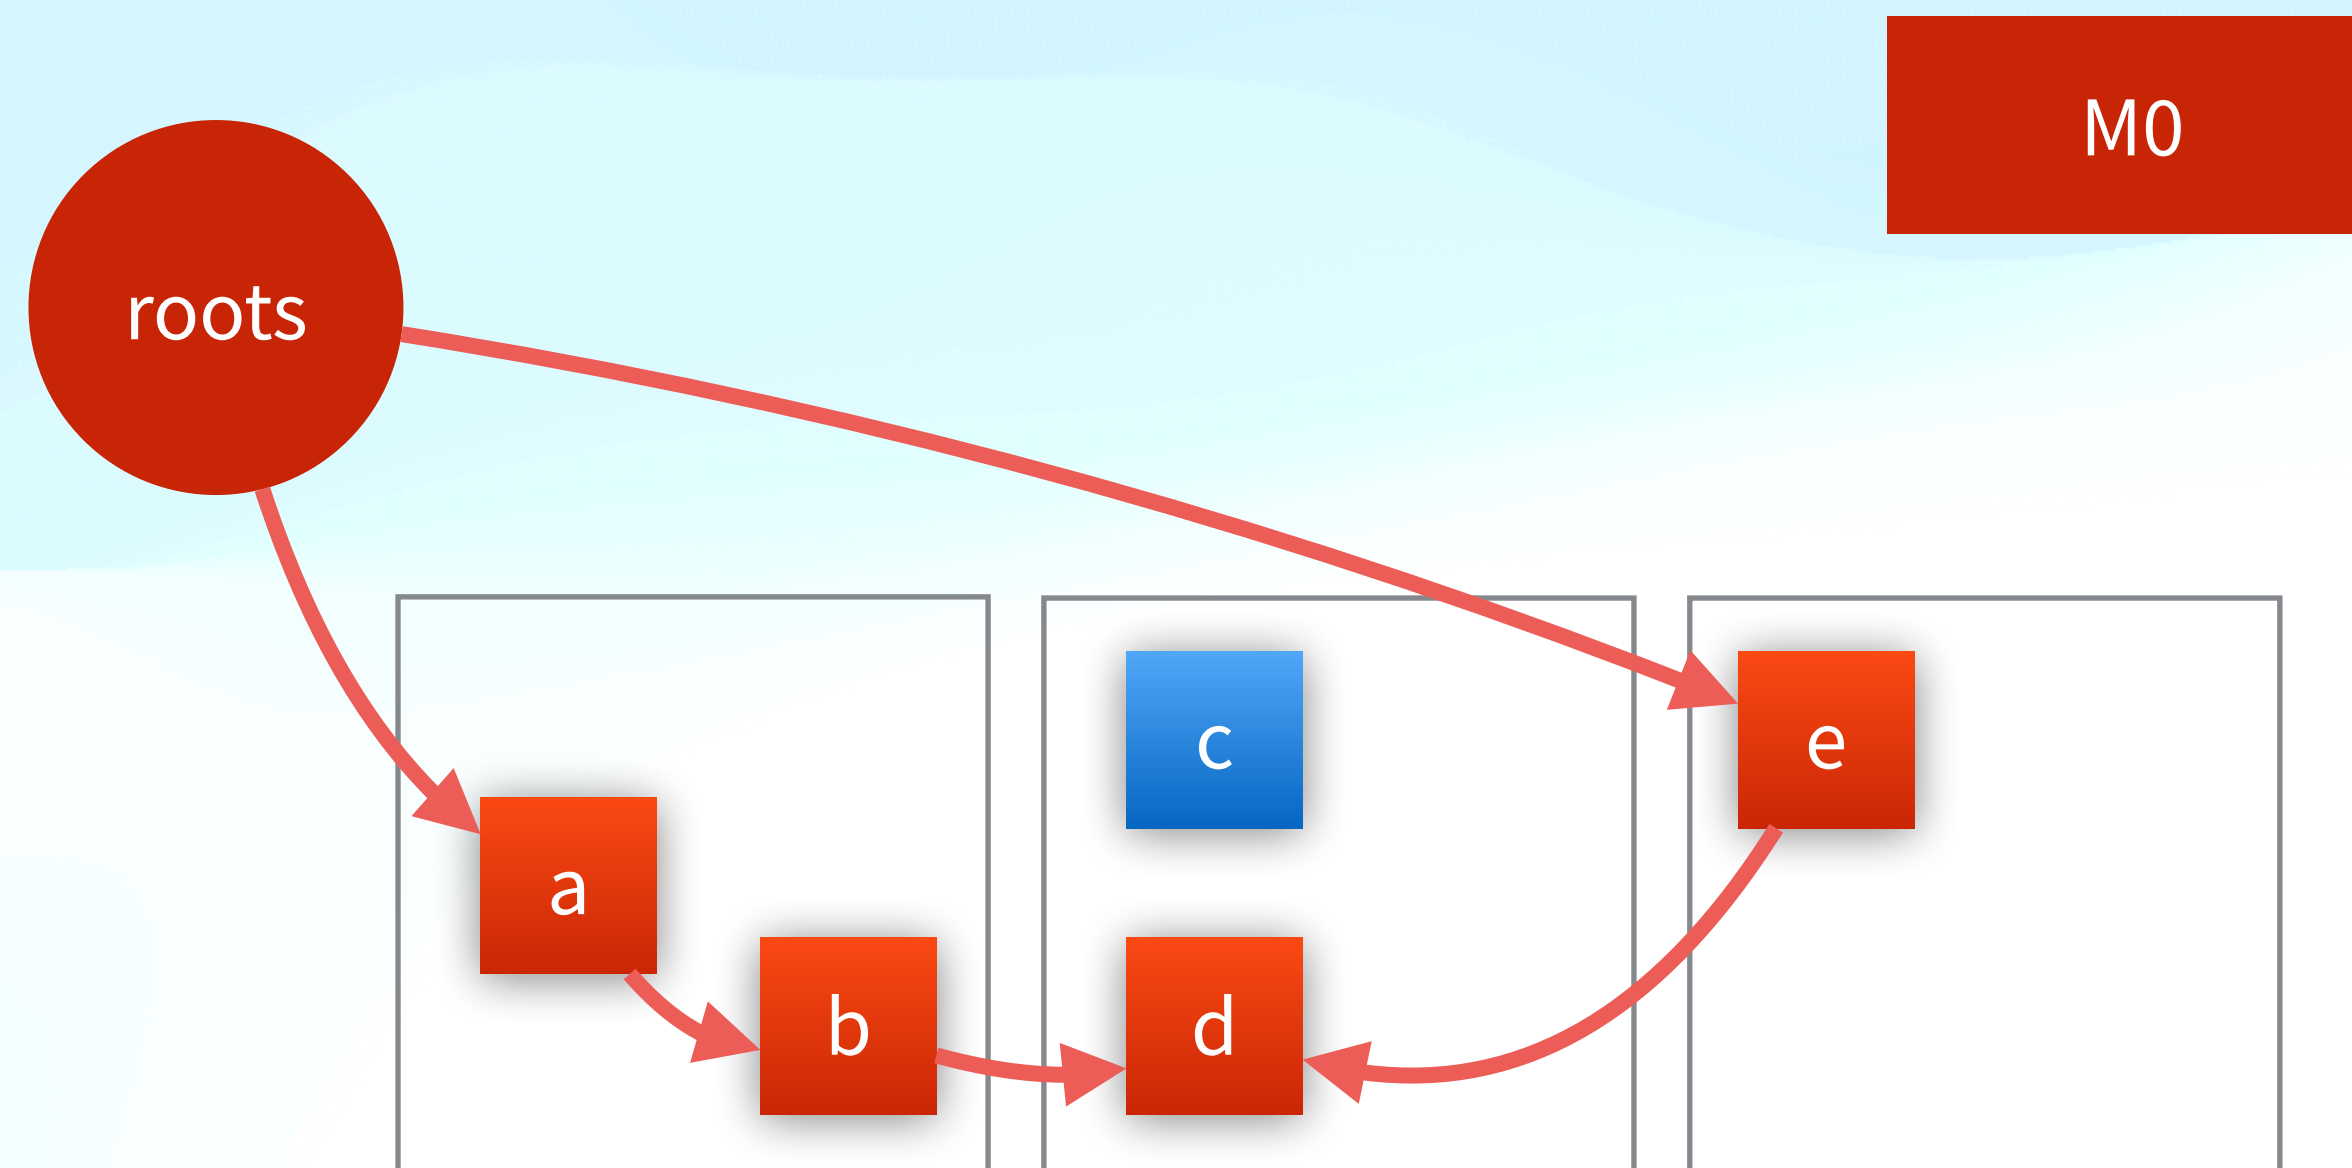


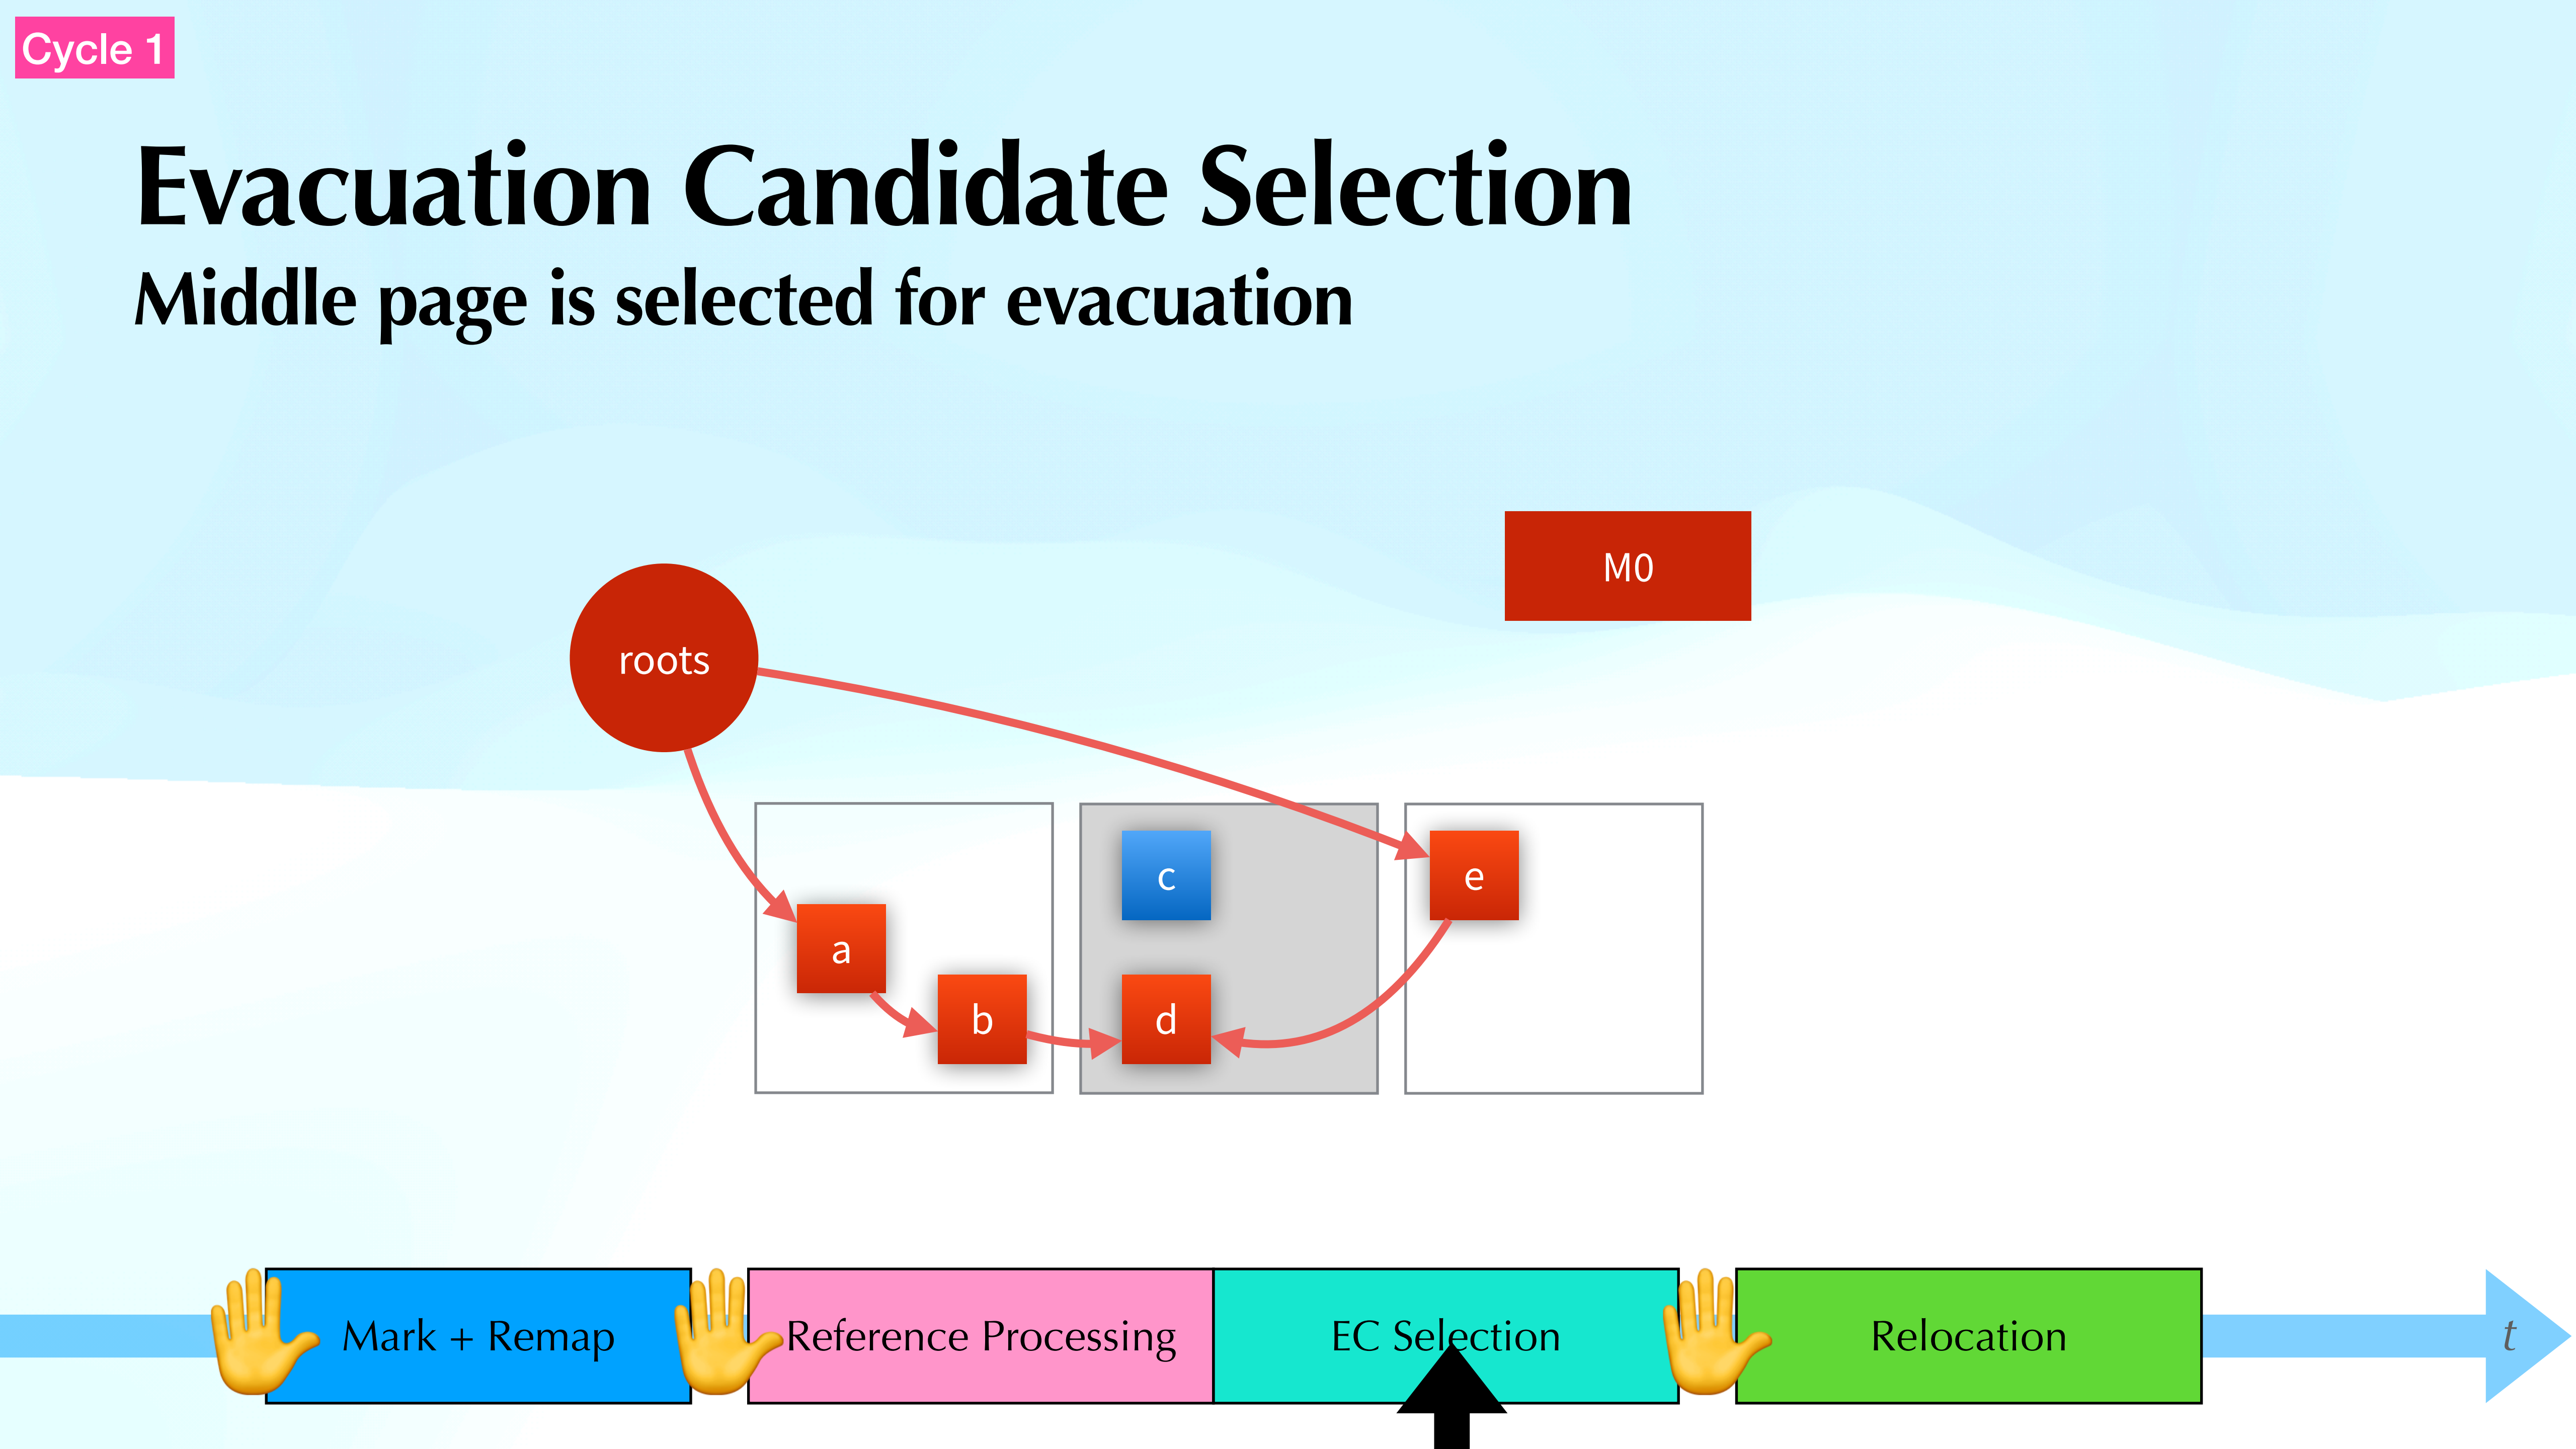


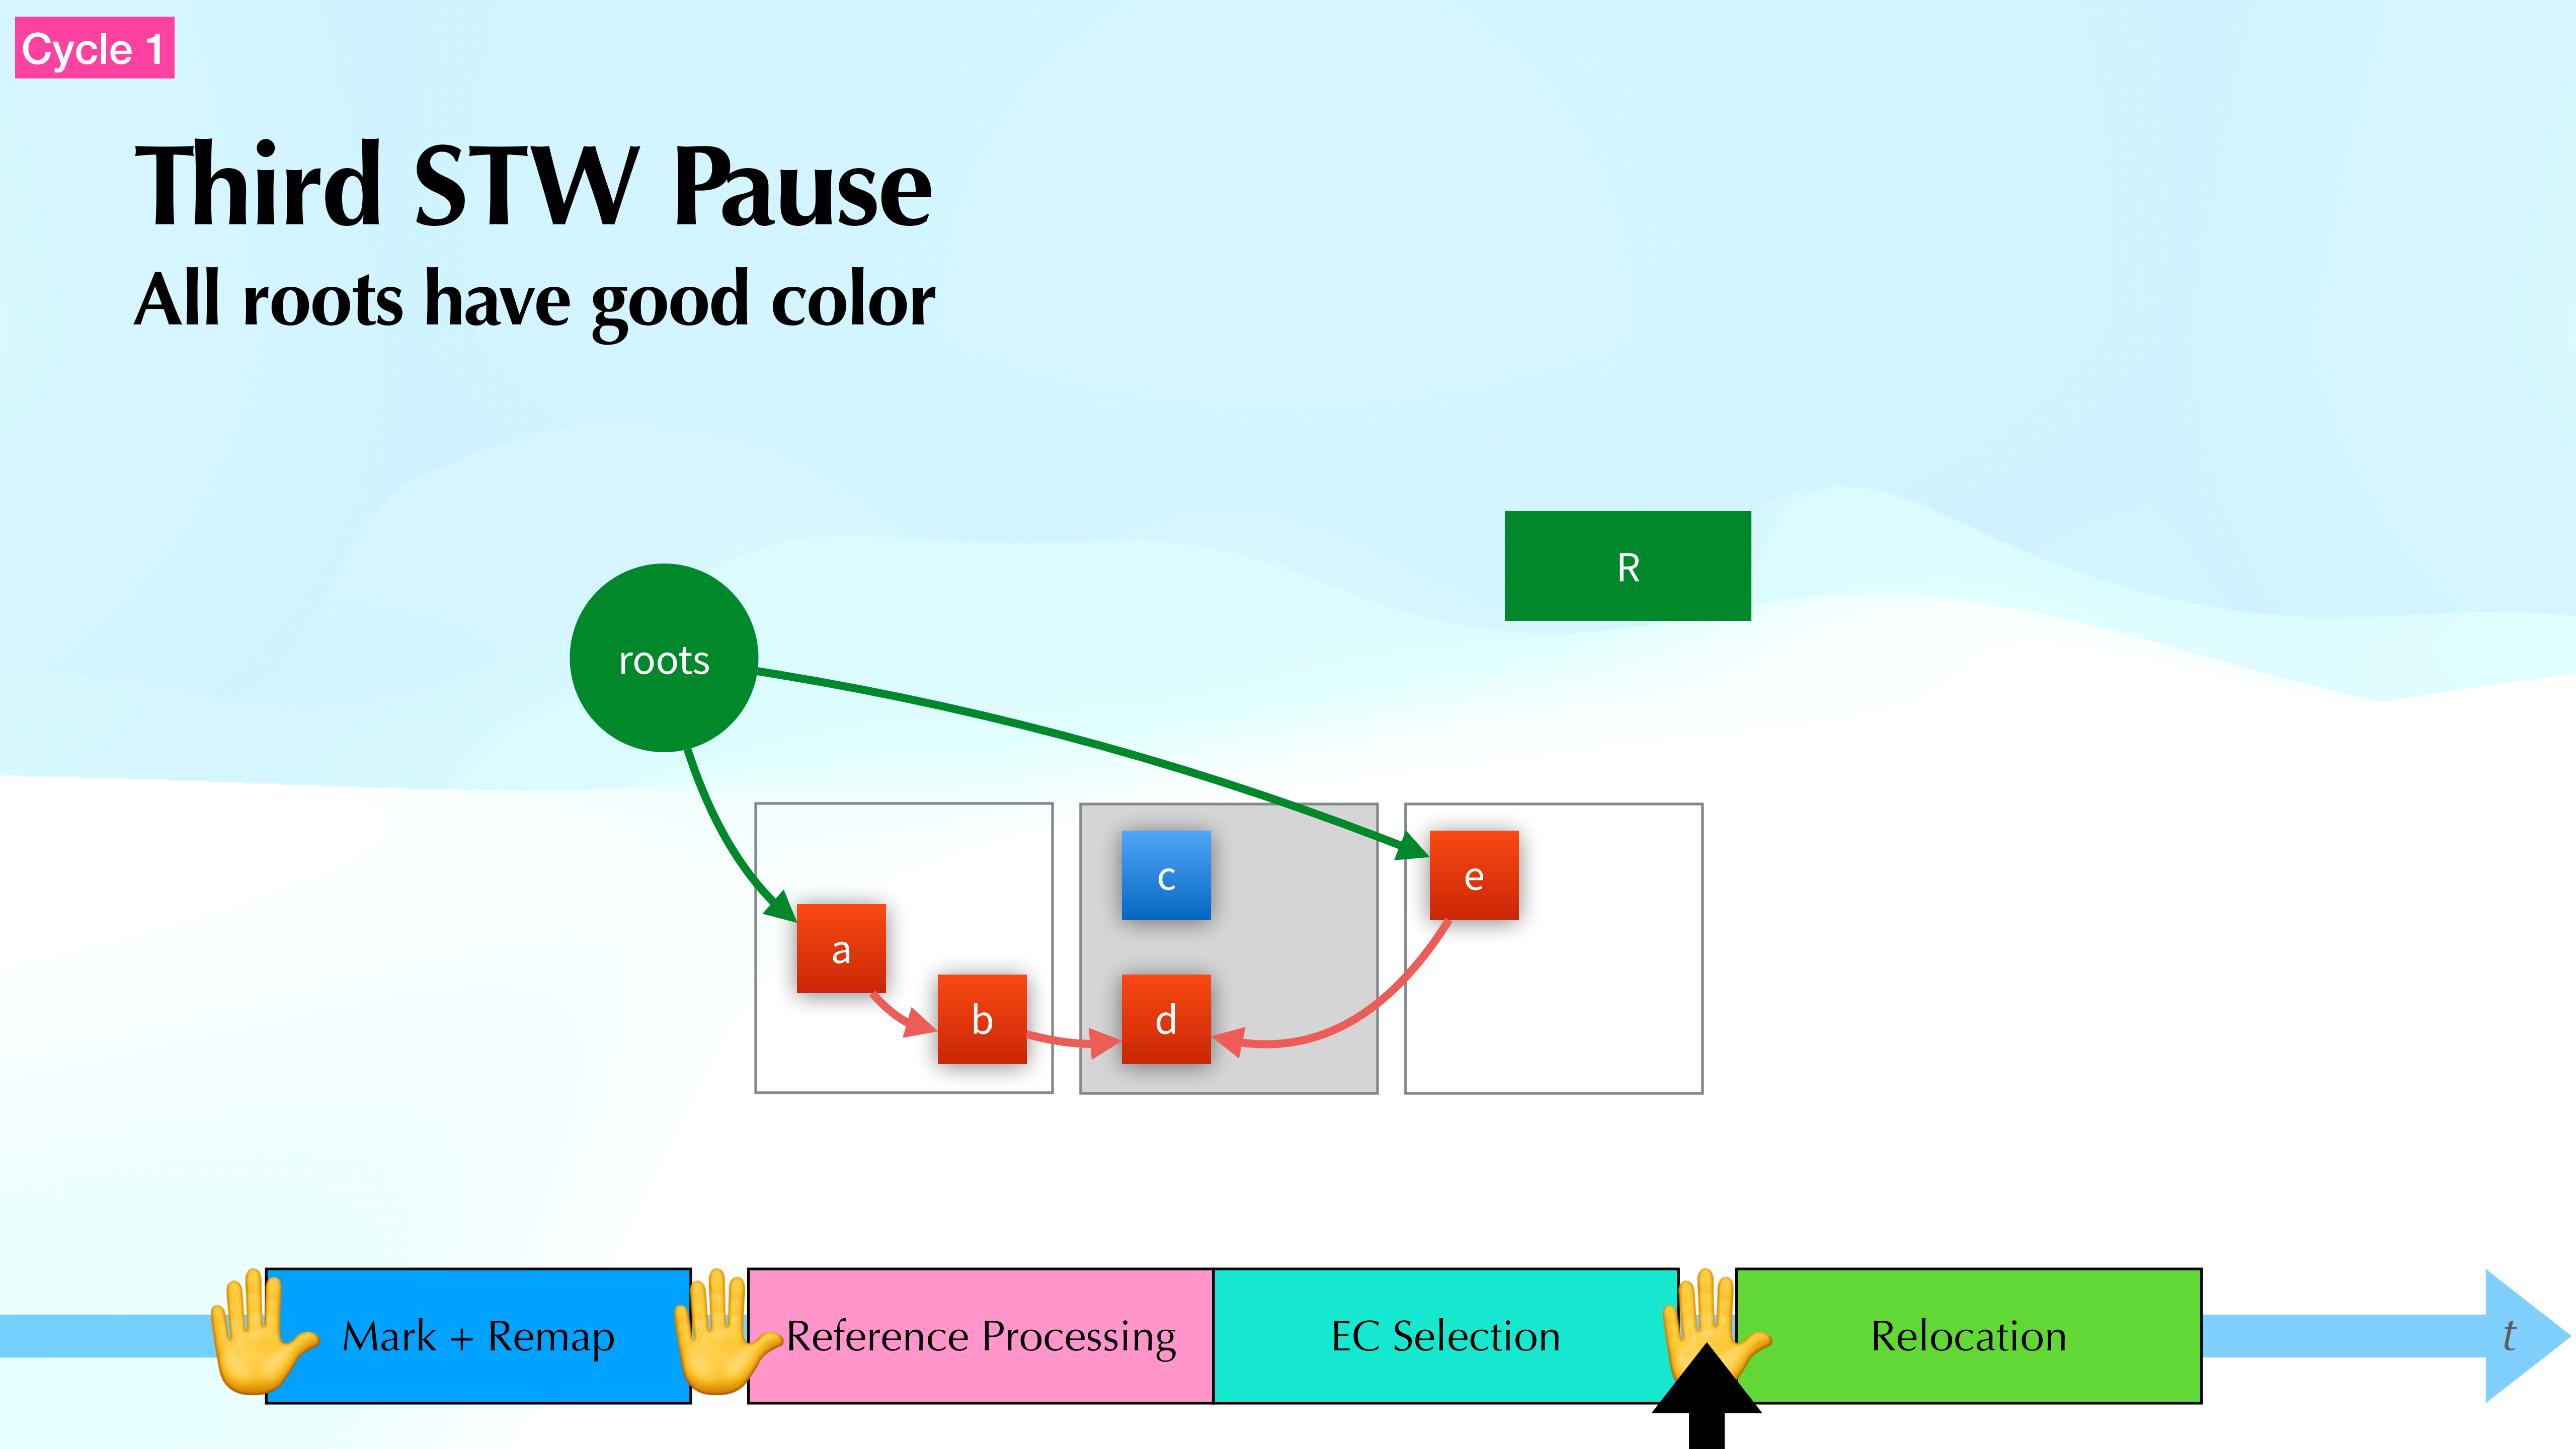


Marking & Remapping

All reachable objects are marked

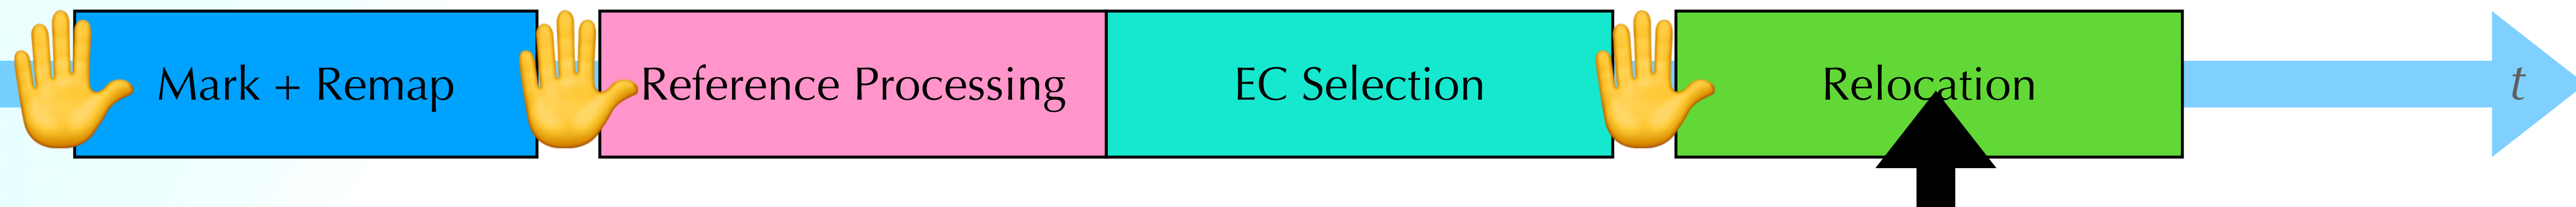



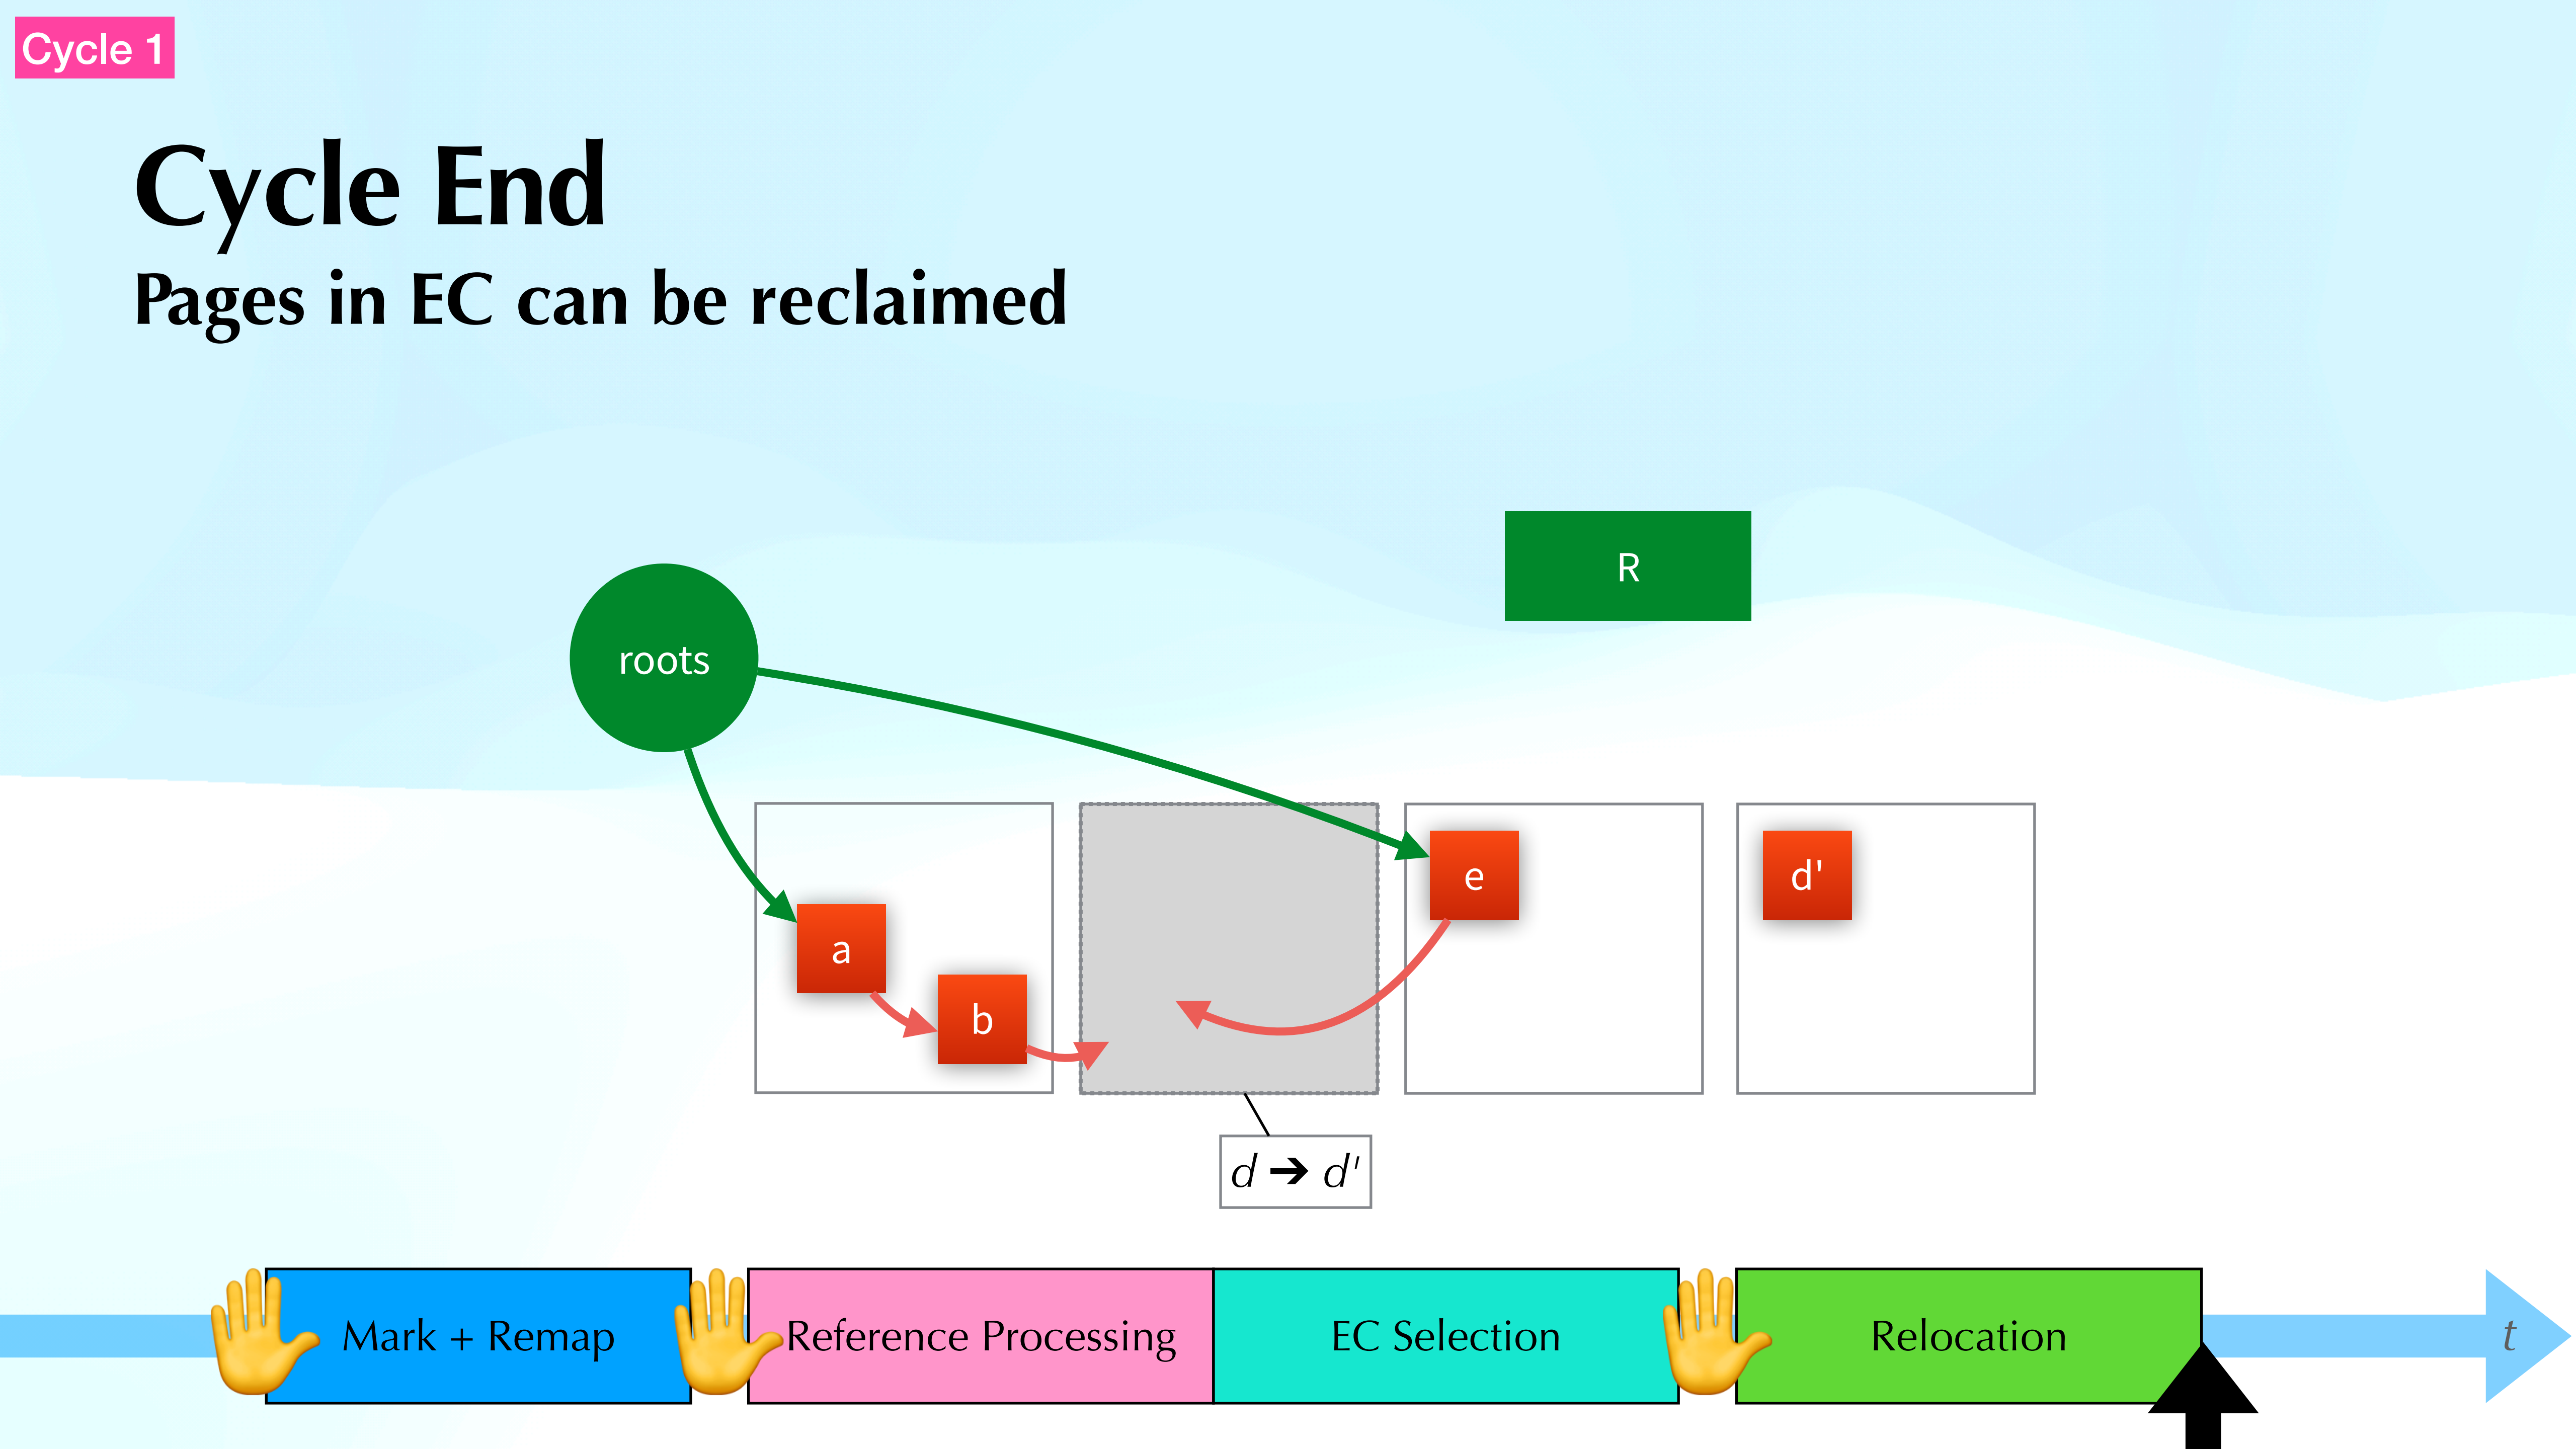


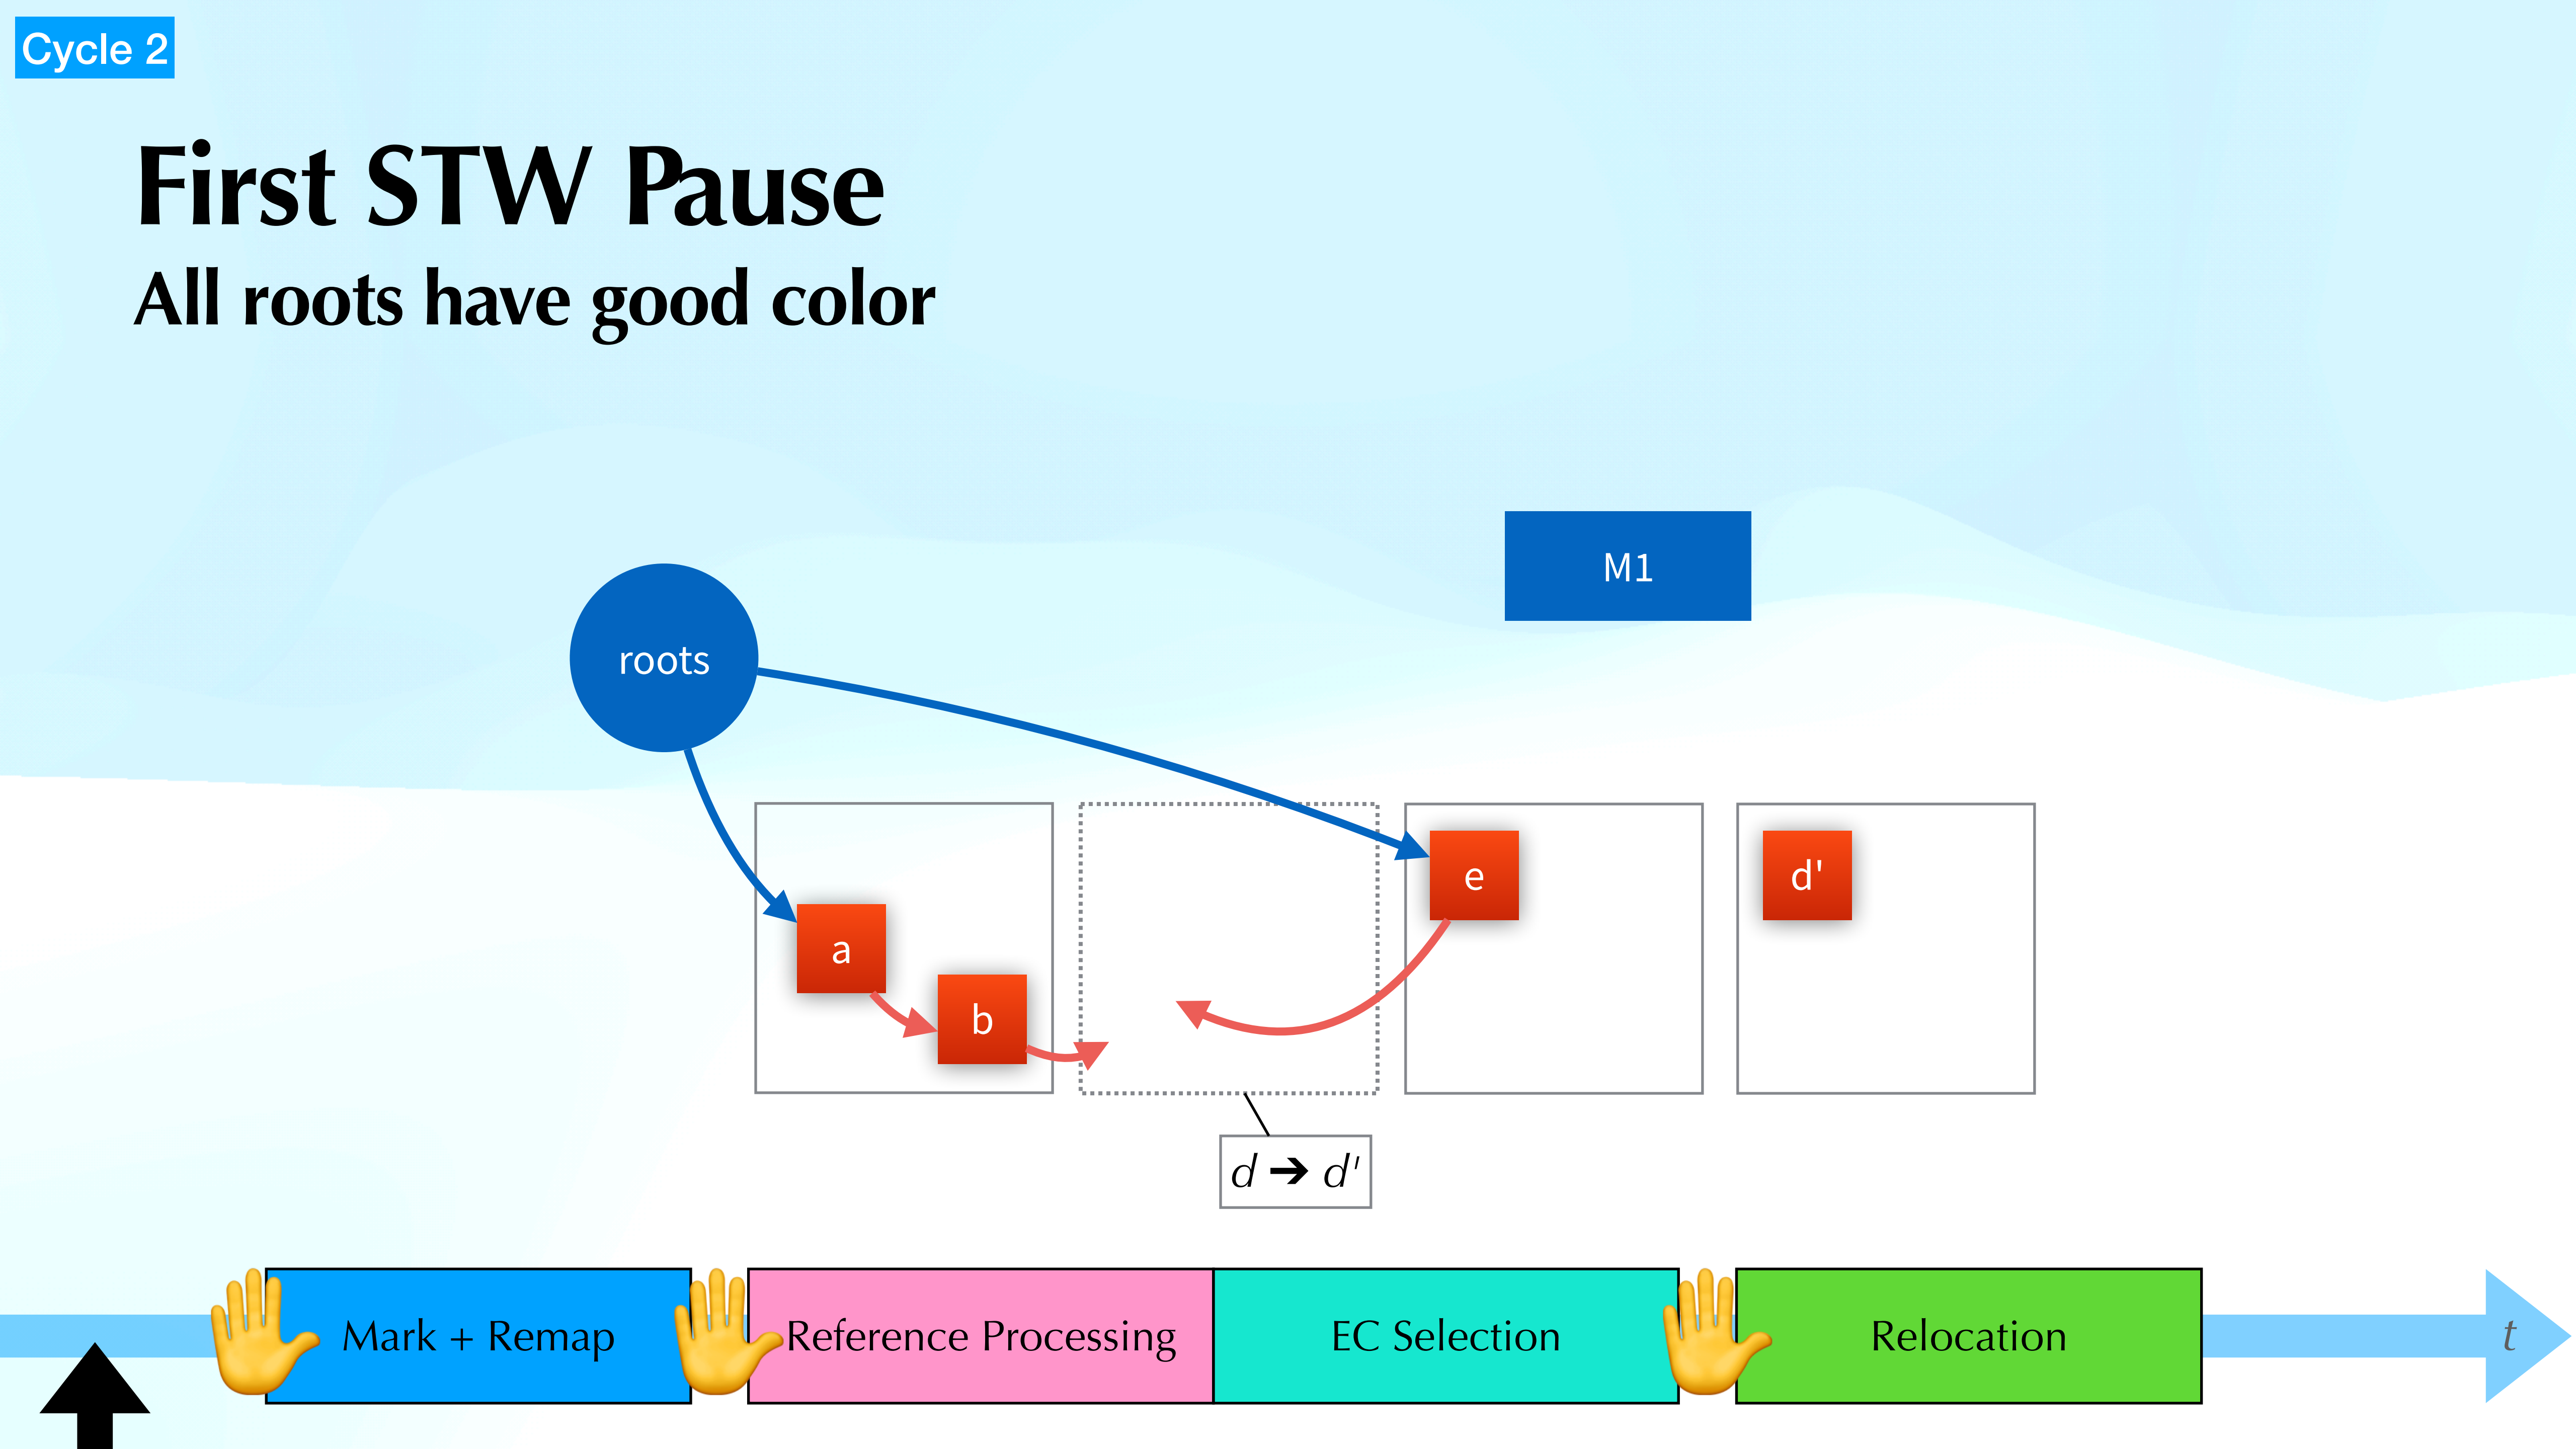


Relocation

EC Pages are evacuated (live objects copied to other pages)

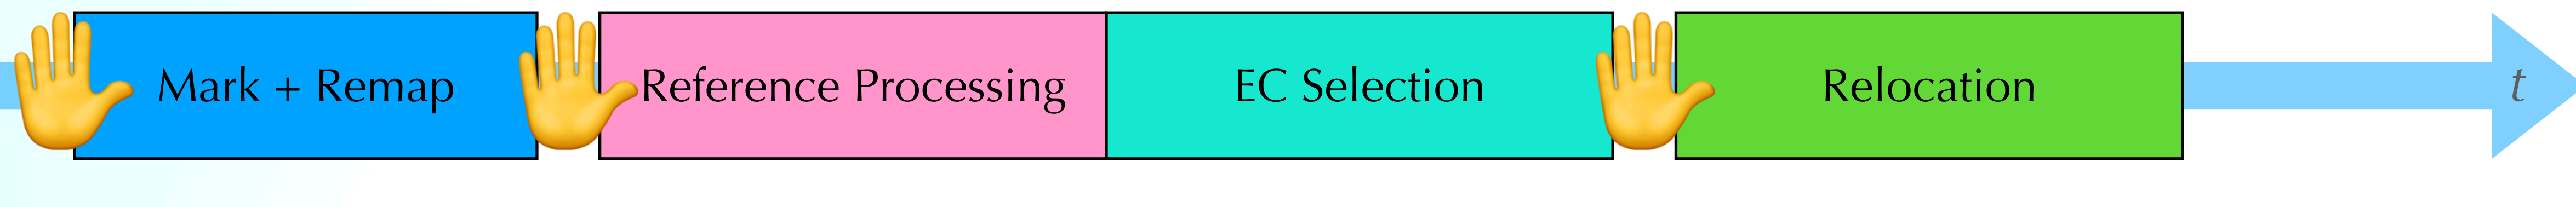







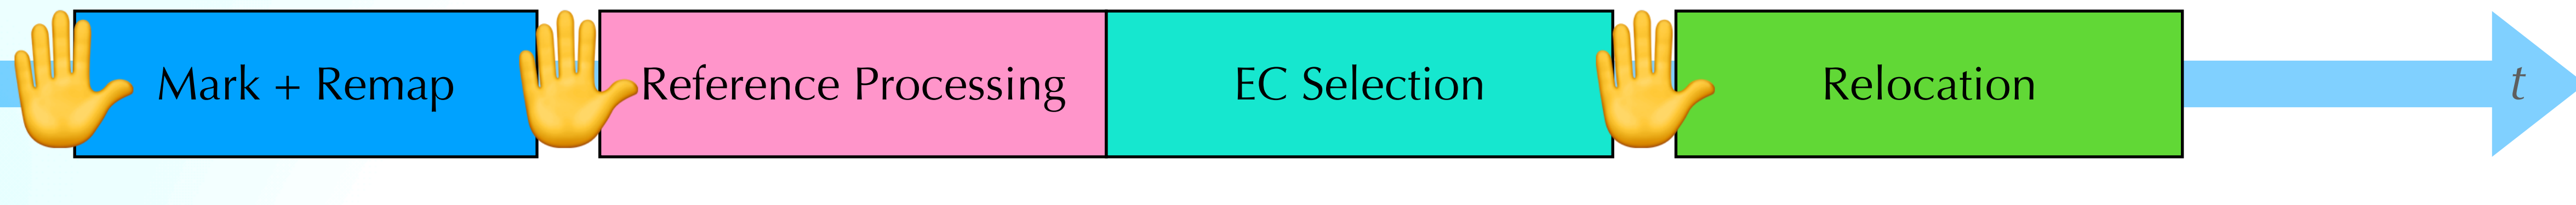
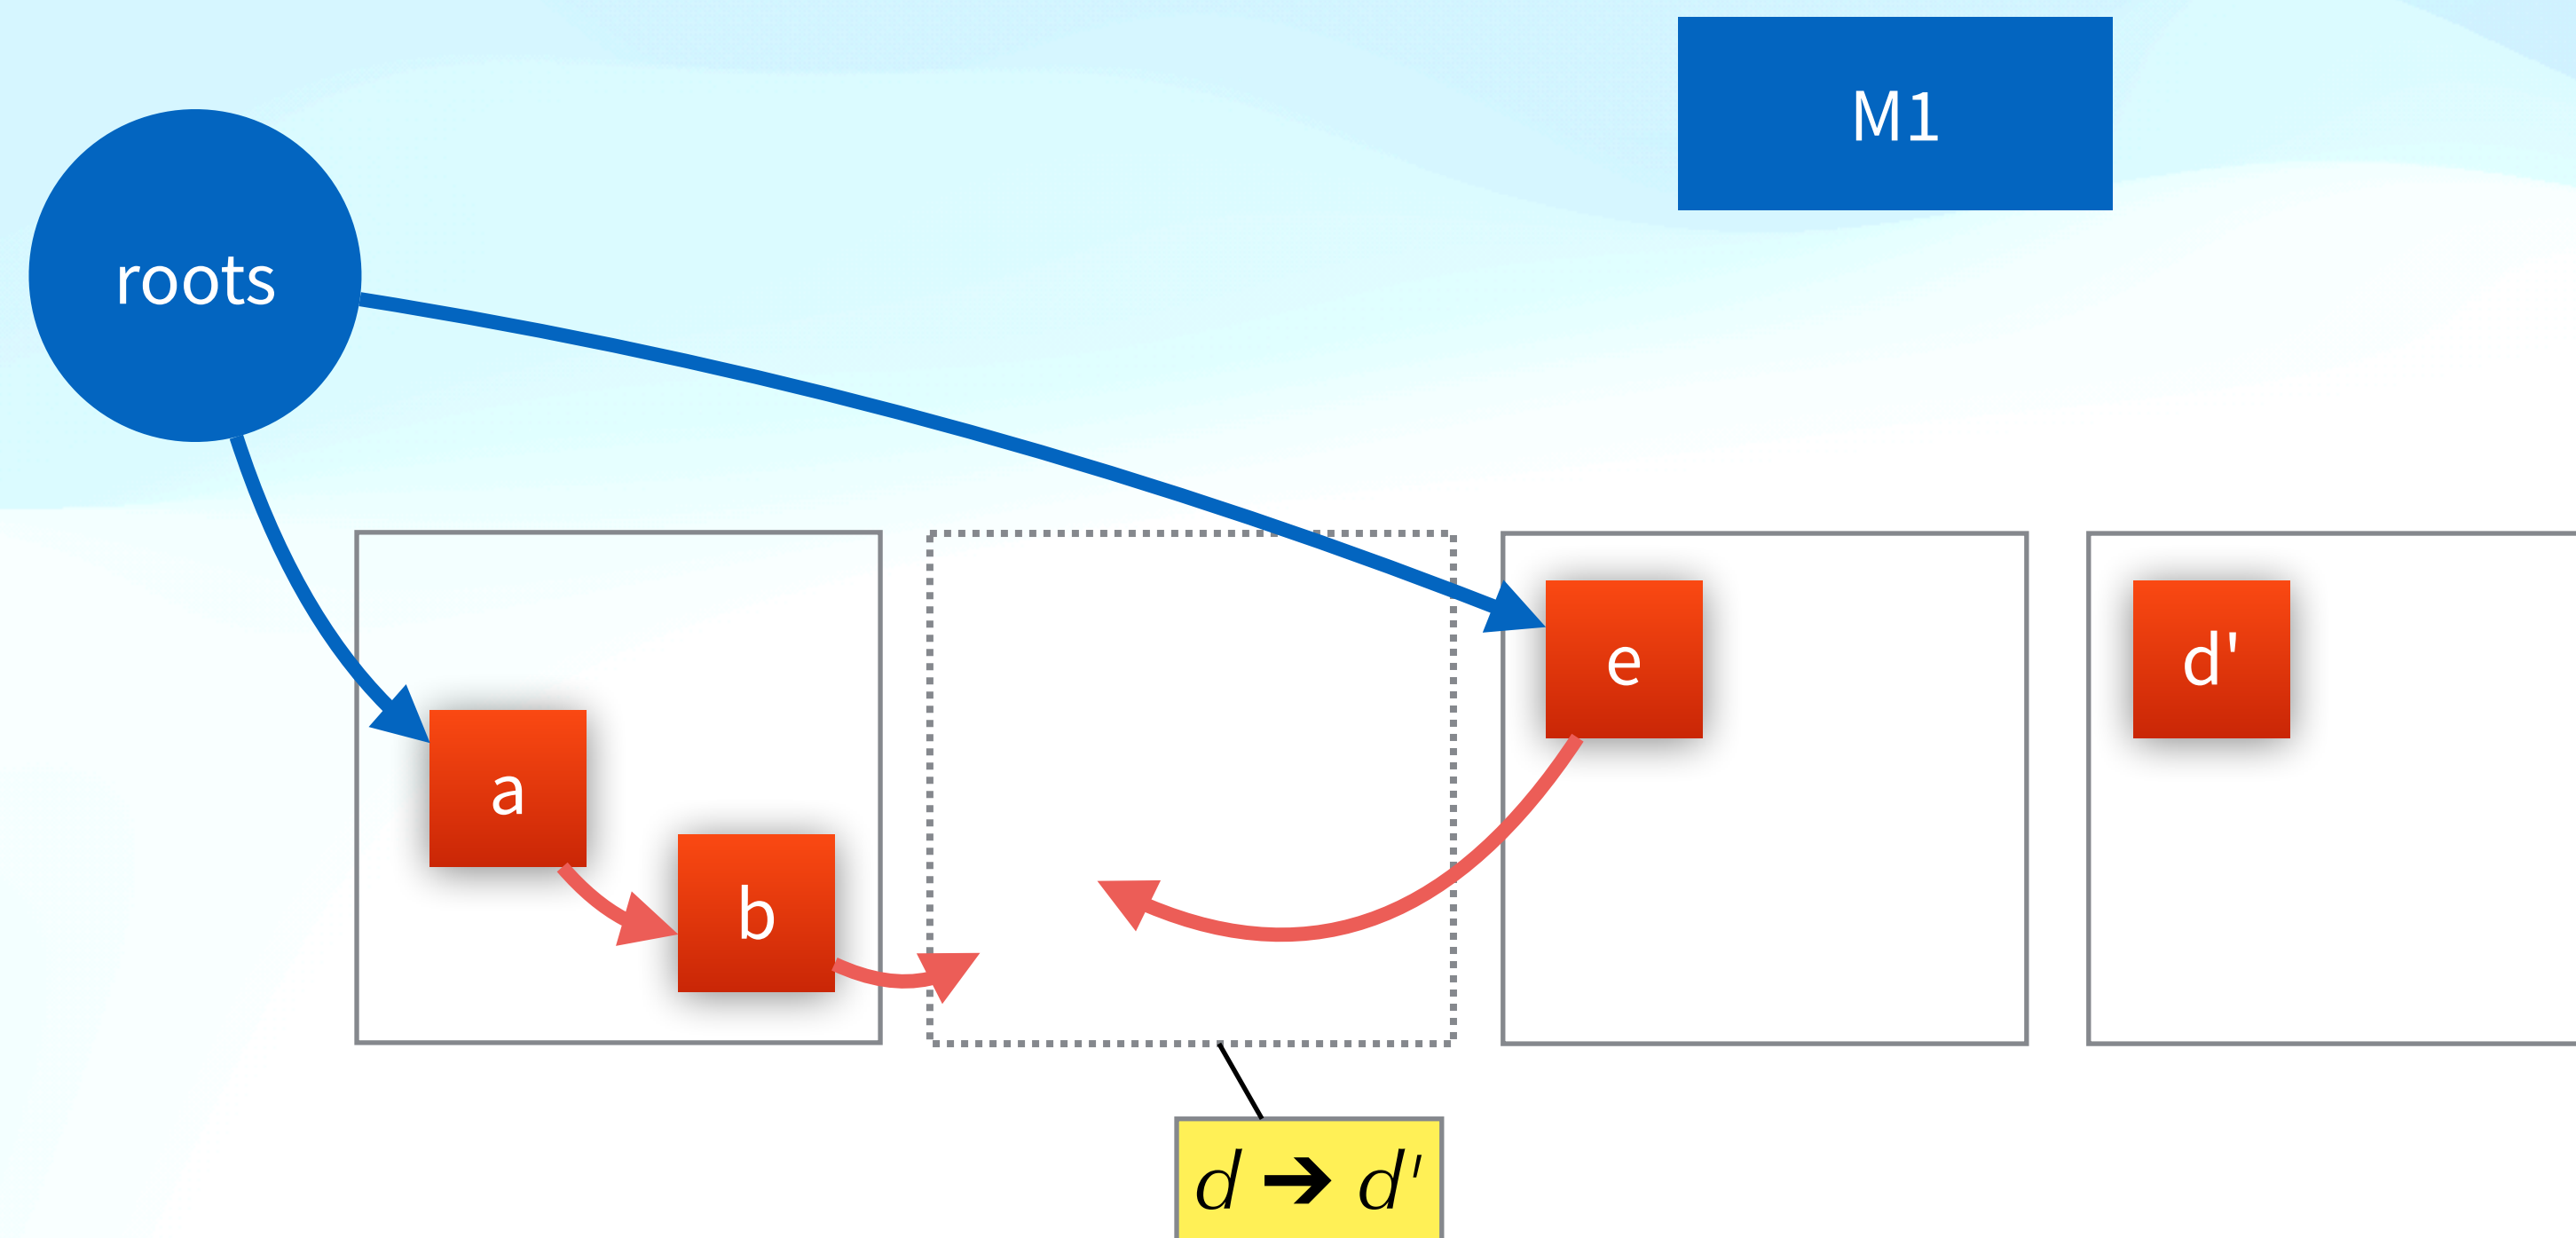
Mutator accesses d via e

Dangling Pointer with Bad Colour — Trapped by Load Barrier





Mutator accesses d via e

Load barrier hits slow path — consults the forwarding table



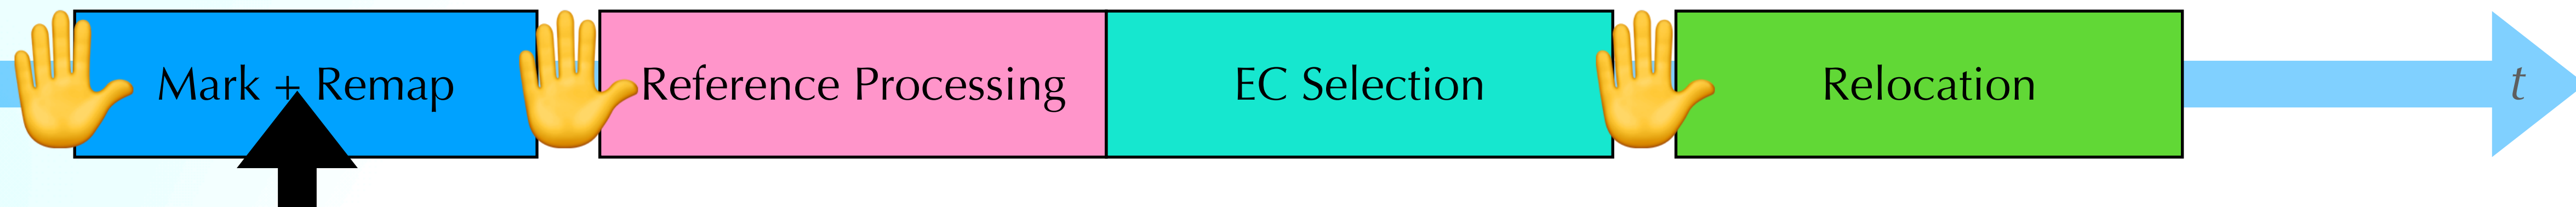

Mutator accesses d via e

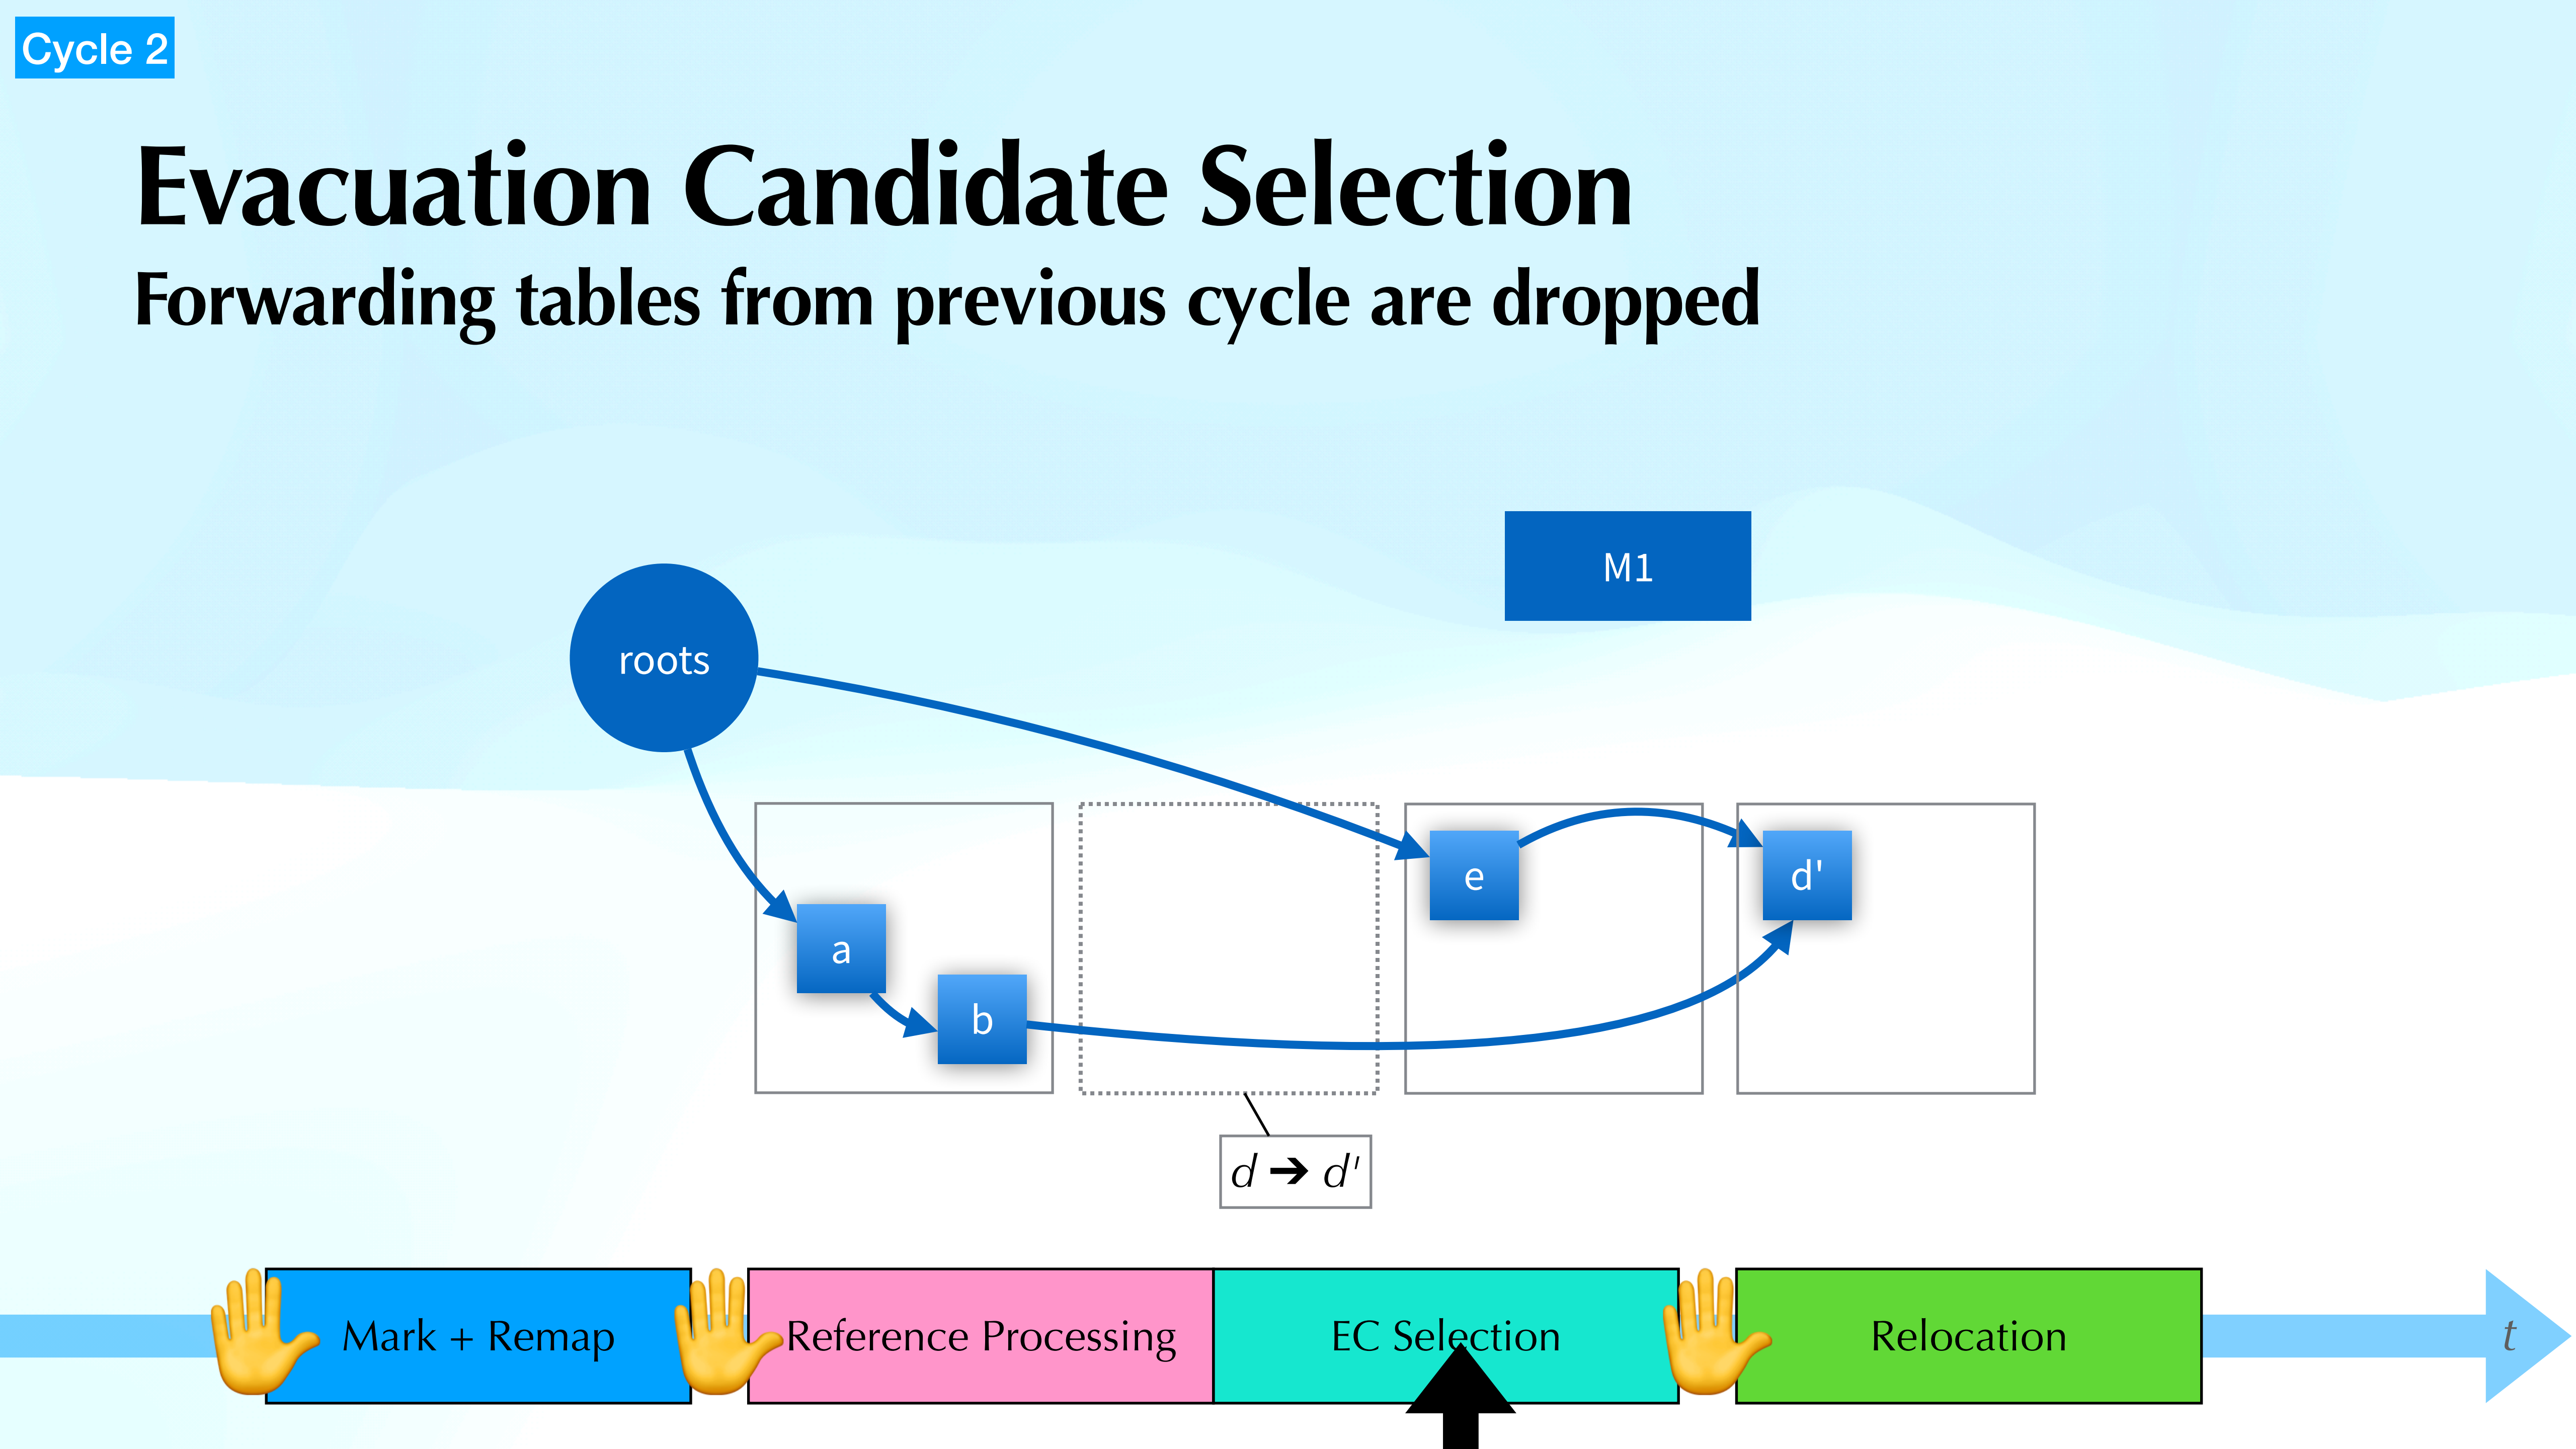
Performs self-healing to avoid future slow paths, and access continues



Marking & Remapping

All outdated pointers are remapped





Summary Slide

- Threads synchronise briefly on what is the good colour
- All pointers are updated to have the good colour once per GC cycle (roots in each STW)
- Dangling pointers will be trapped by load barriers
- Evacuated pages can be dropped immediately
- Forwarding tables can be dropped after following marking phase

Aren't load barriers too expensive?

- Nah; Compressed OOPS
- They can be — it depends on the load-barrier design

"If you screw up and make a load barrier that is more than 2 instructions, you've made your own mess"

ZGC Load Barrier x86_64 Instructions

```
movq rax, 0x10(rbx)
testq rax, 0x20(r15)
jnz slow_path
```


- ZGC used multi mapping to make coloured addresses valid



Extending ZGC to improve locality

Yang, Österlund, and Wrigstad (PLDI 2020)

- Good spatial locality can hide memory latency
- Managed languages tend to abstract memory
 - Fewer bugs
 - Harder to optimise




- **RQ:** how can we help programmers get good locality in managed languages?

Extending ZGC to improve locality

Yang, Österlund, and Wrigstad (PLDI 2020)

- Good spatial locality can hide memory latency
- Managed languages tend to abstract memory
 - Fewer bugs
 - Harder to optimise

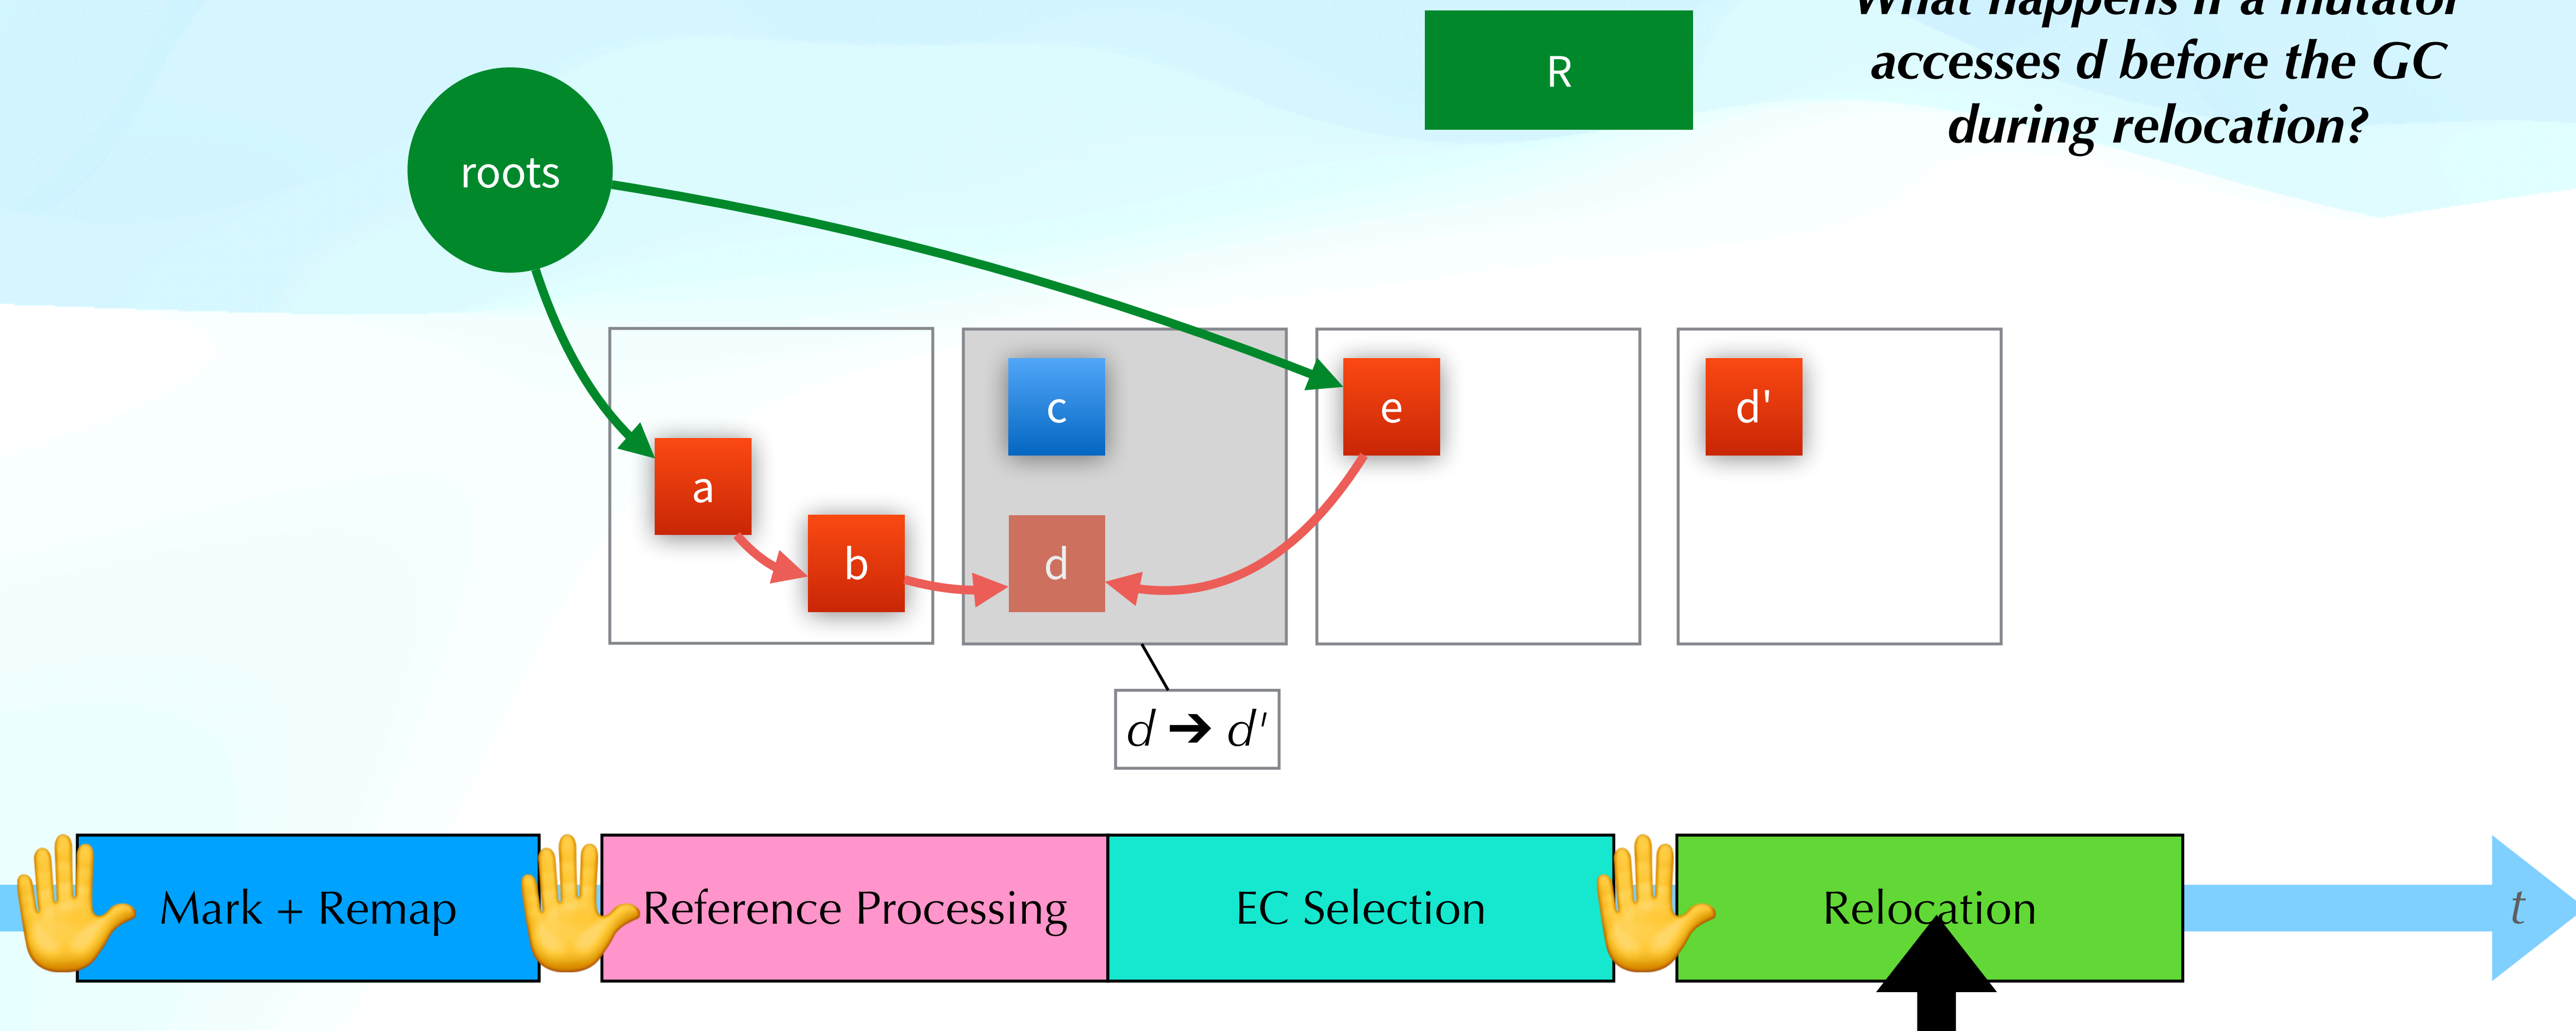


- **RQ:** how can we help programmers get good locality in managed languages?

Extending ZGC to improve locality

Yang, Österlund, and Wrigstad (PLDI 2020)

What happens if a mutator accesses d before the GC during relocation?



Design mismatches and mitigations

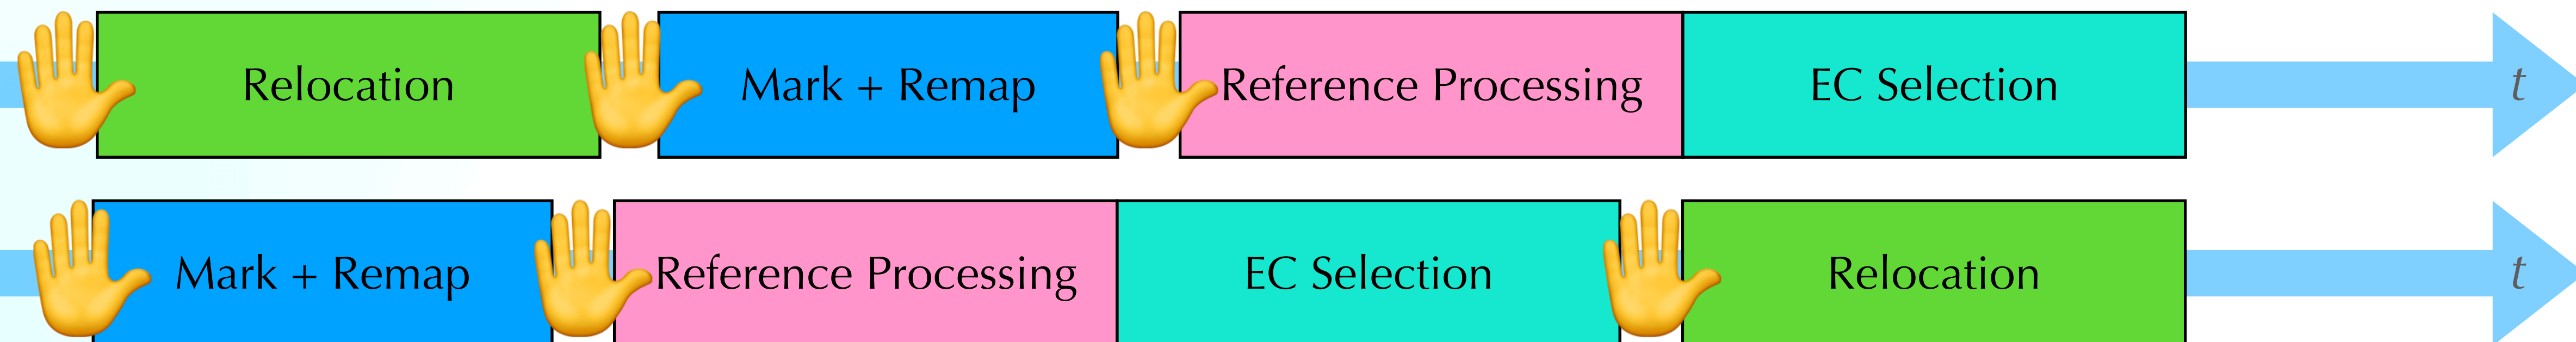
- In ZGC, relocation starts immediately after EC selection
 - Thus, mutators are competing with GC threads to move objects
 - **Question:** how can we help the mutators win that race?
- Only objects on pages in EC are movable
 - **Question:** how do we make more objects movable?



Design mismatches and mitigations

"Hacks": **change phase order in the GC cycle**; **add all pages to EC**

- In ZGC, relocation starts immediately after EC selection
 - Thus, mutators are competing with GC threads to move objects
 - **Question**: how can we help the mutators win that race?
- Only objects on pages in EC are movable
 - **Question**: how do we make more objects movable?



Config ID	Hotness	ColdPage	ColdConfidence	RelocateAllPages	LazyRelocate
1	0	0	0.0	0	0
2	0	0	0.0	0	1
3	0	0	0.0	1	0
4	0	0	0.0	1	1
5	1	0	0.0	0	0
6	1	0	0.5	0	0
7	1	0	1.0	0	0
8	1	0	0.0	0	1
9	1	0	0.5	0	1
10	1	0	1.0	0	1
11	1	1	0.0	0	0
12	1	1	0.5	0	0
13	1	1	1.0	0	0
14	1	1	0.0	0	1
15	1	1	0.5	0	1
16	1	1	1.0	0	1
17	1	1	0.0	1	0
18	1	1	0.0	1	1

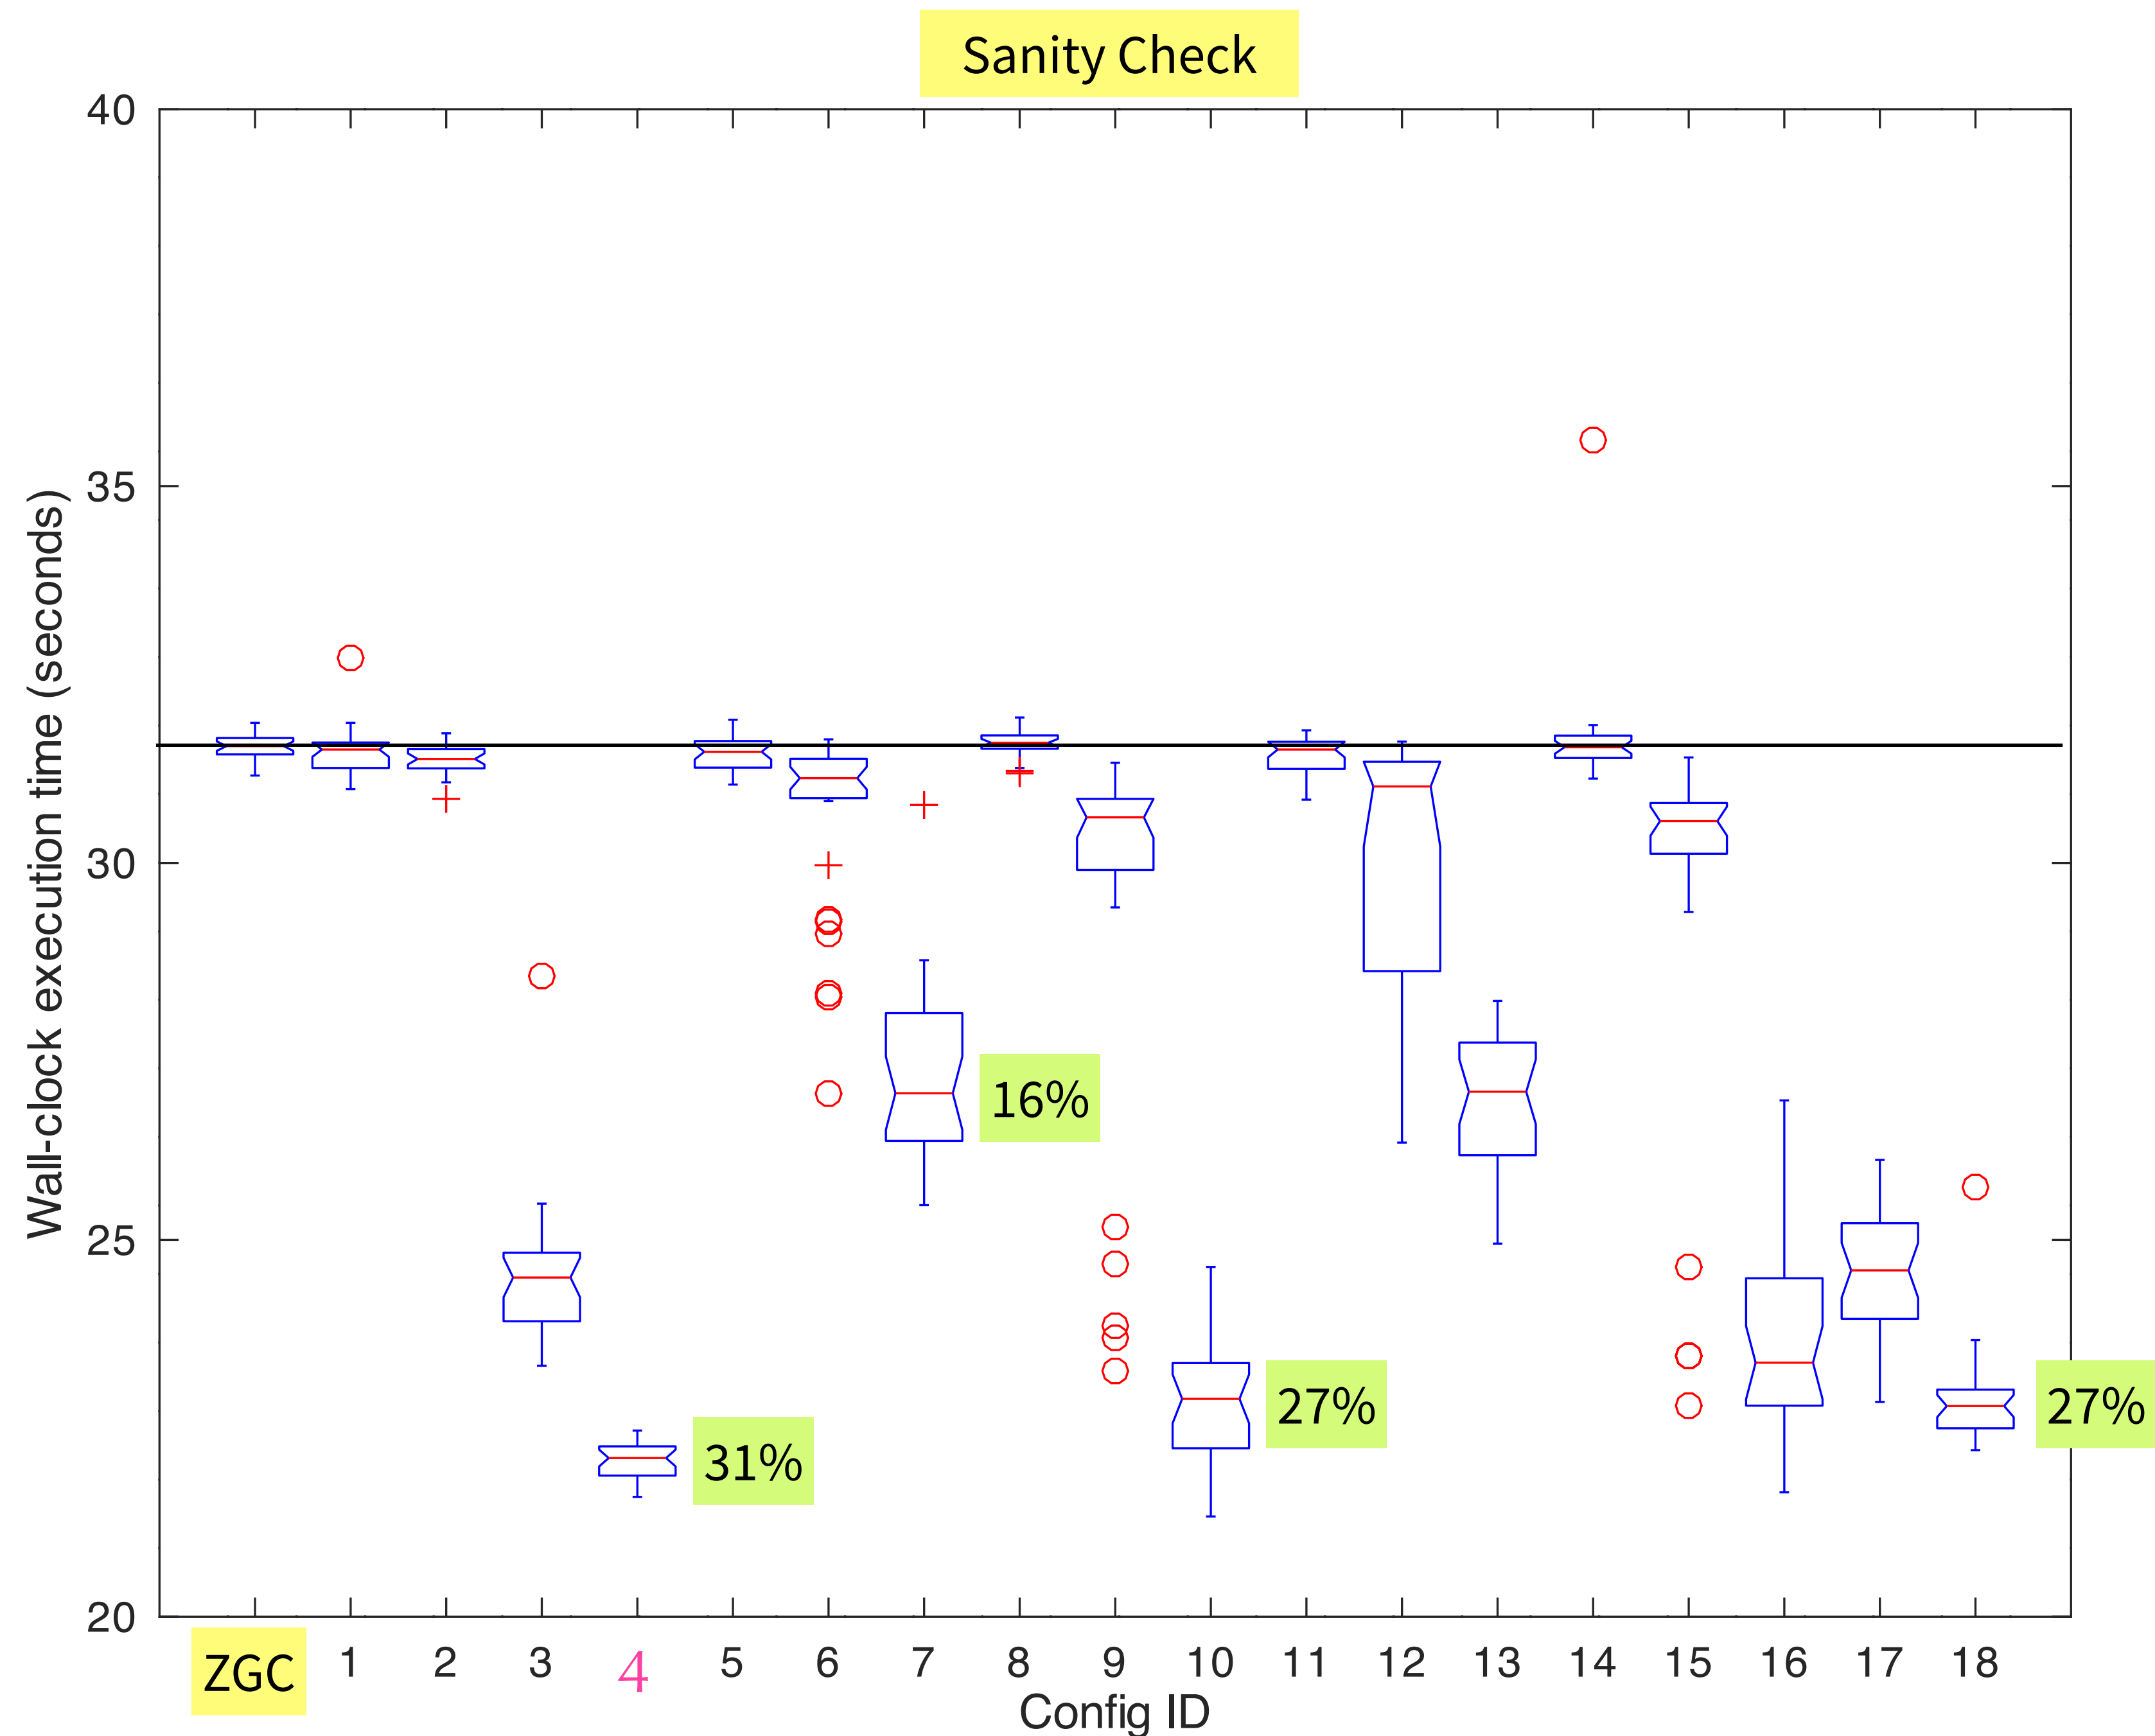


4

Sanity Check Result



(Expect better performance)

HCSGC





Why?

Cache performance




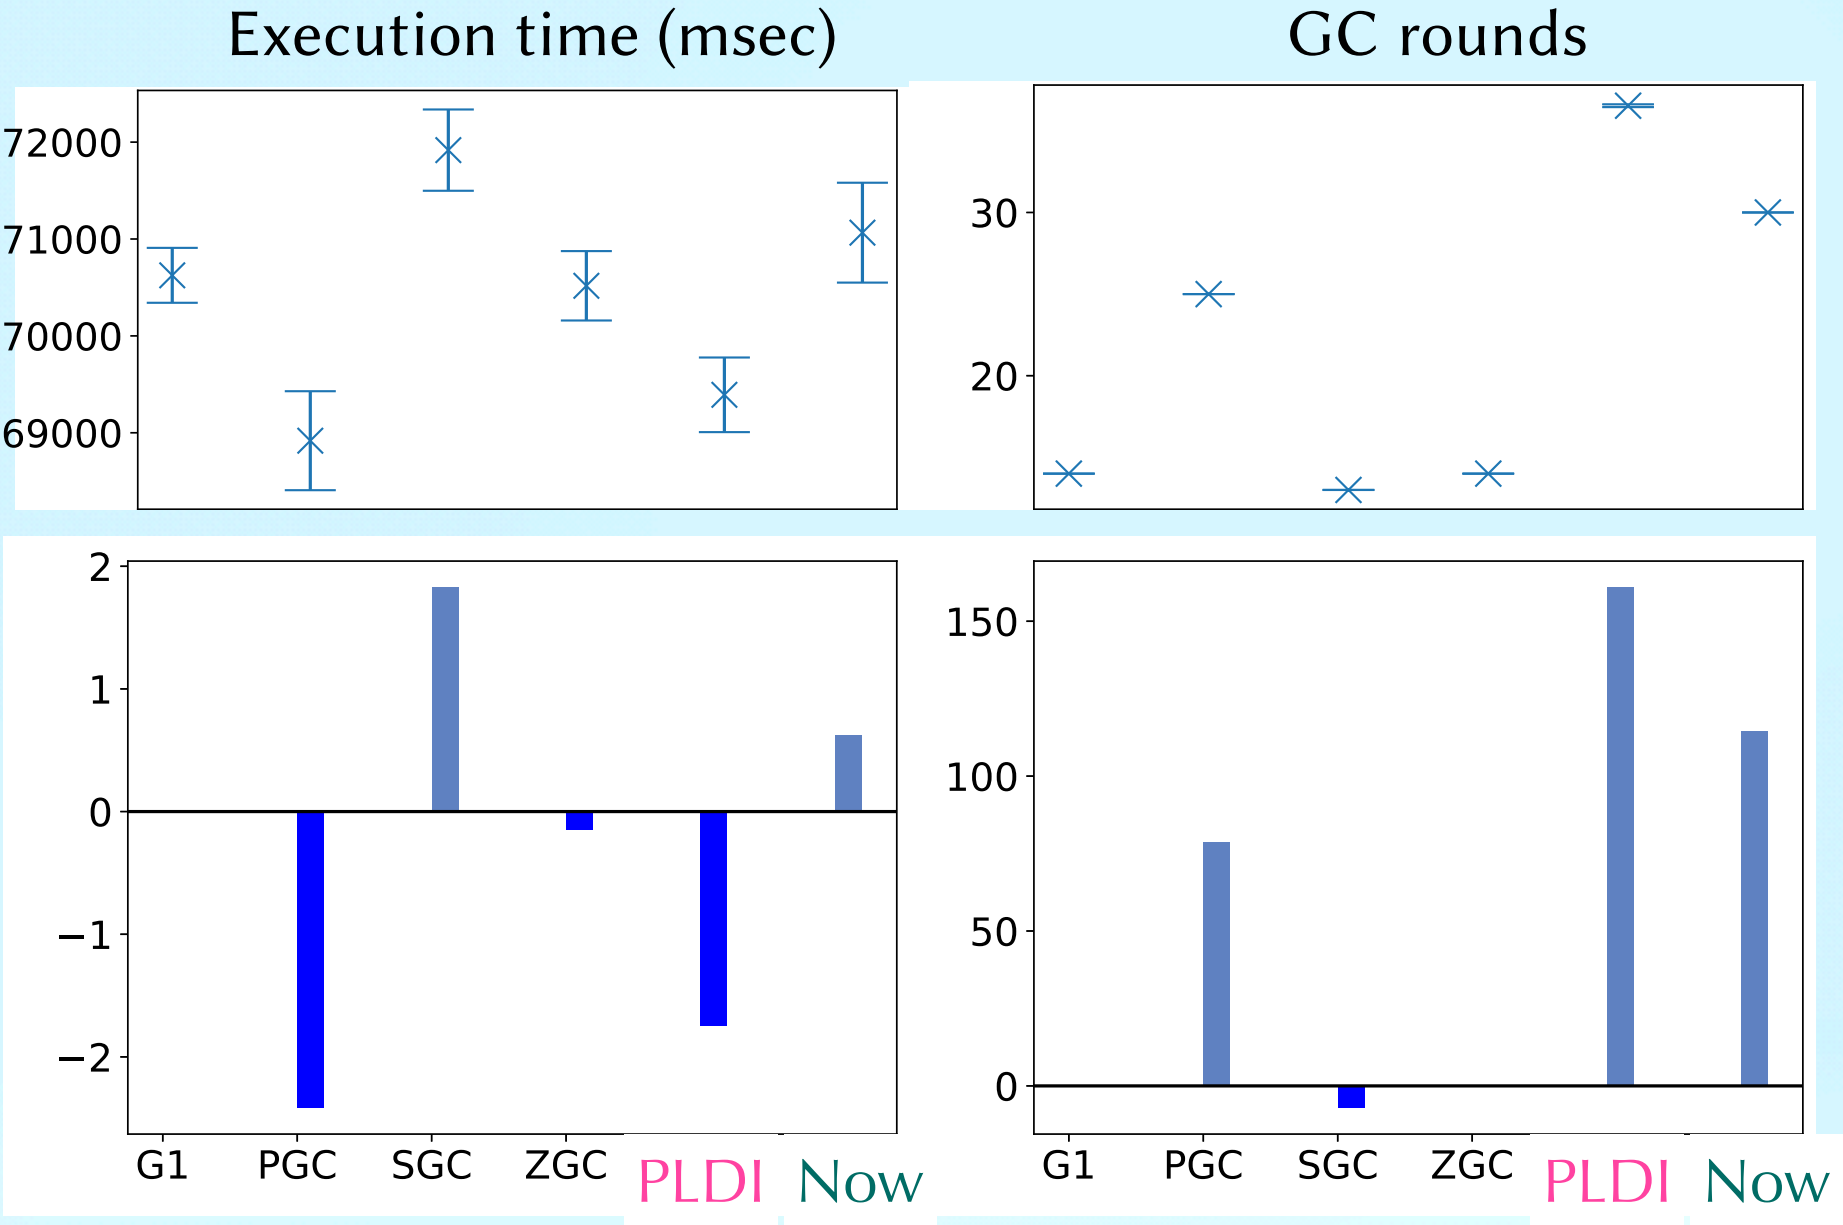


JGraphT Results

Best case for our technique



Continuation
of this work



tradebeans

h2

Speculative Optimisation

- We are letting the mutator copy objects hoping that will improve performance
- Filtering is key — don't want to copy objects...
 - ...that are too large (almost no point)
 - ...that are already well-placed
 - ...that won't be accessed again
 - ...etc.
- **RQ:** how can we make that filtering (and efficiently)?

Generational ZGC

Generational Garbage Collectors

For programs that satisfy the generational hypothesis

- Divide the heap into multiple (two) spaces (young, old) managed separately
- Objects in young space likely to die — high ROI on GC in young space
- Objects in old space less likely to die — lower ROI on GC in old space
- Reduced Effort to Collect Garbage
 - Withstand higher allocation rates
 - Lower heap headroom
 - Lower CPU usage

Heap



The diagram illustrates the structure of the Java Heap memory space. It is represented as a large horizontal rectangle divided into two equal-width vertical sections. The left section is a light blue color and is labeled 'Young Generation'. The right section is a darker blue color and is labeled 'Old Generation'. A black curly brace is positioned above the top edge of the rectangle, spanning its entire width and centered under the 'Heap' title.

Young Generation

Old Generation

ZGC Heap

Y	Y	Y	O	Y	O	O	O
Y	O	O	O			Y	O
	O	Y	O	Y			
	Y	O		Y	O	O	O

ZGC Heap



The diagram illustrates the ZGC (Zing Garbage Collection) heap structure. It consists of two main memory spaces: the Young Generation and the Old Generation. The Young Generation is represented by a light blue rectangle on the left, and the Old Generation is represented by a dark teal rectangle on the right. A horizontal curly brace spans the top of both rectangles, with the label 'ZGC Heap' centered above it.

Young Generation

Old Generation

ZGC Heap

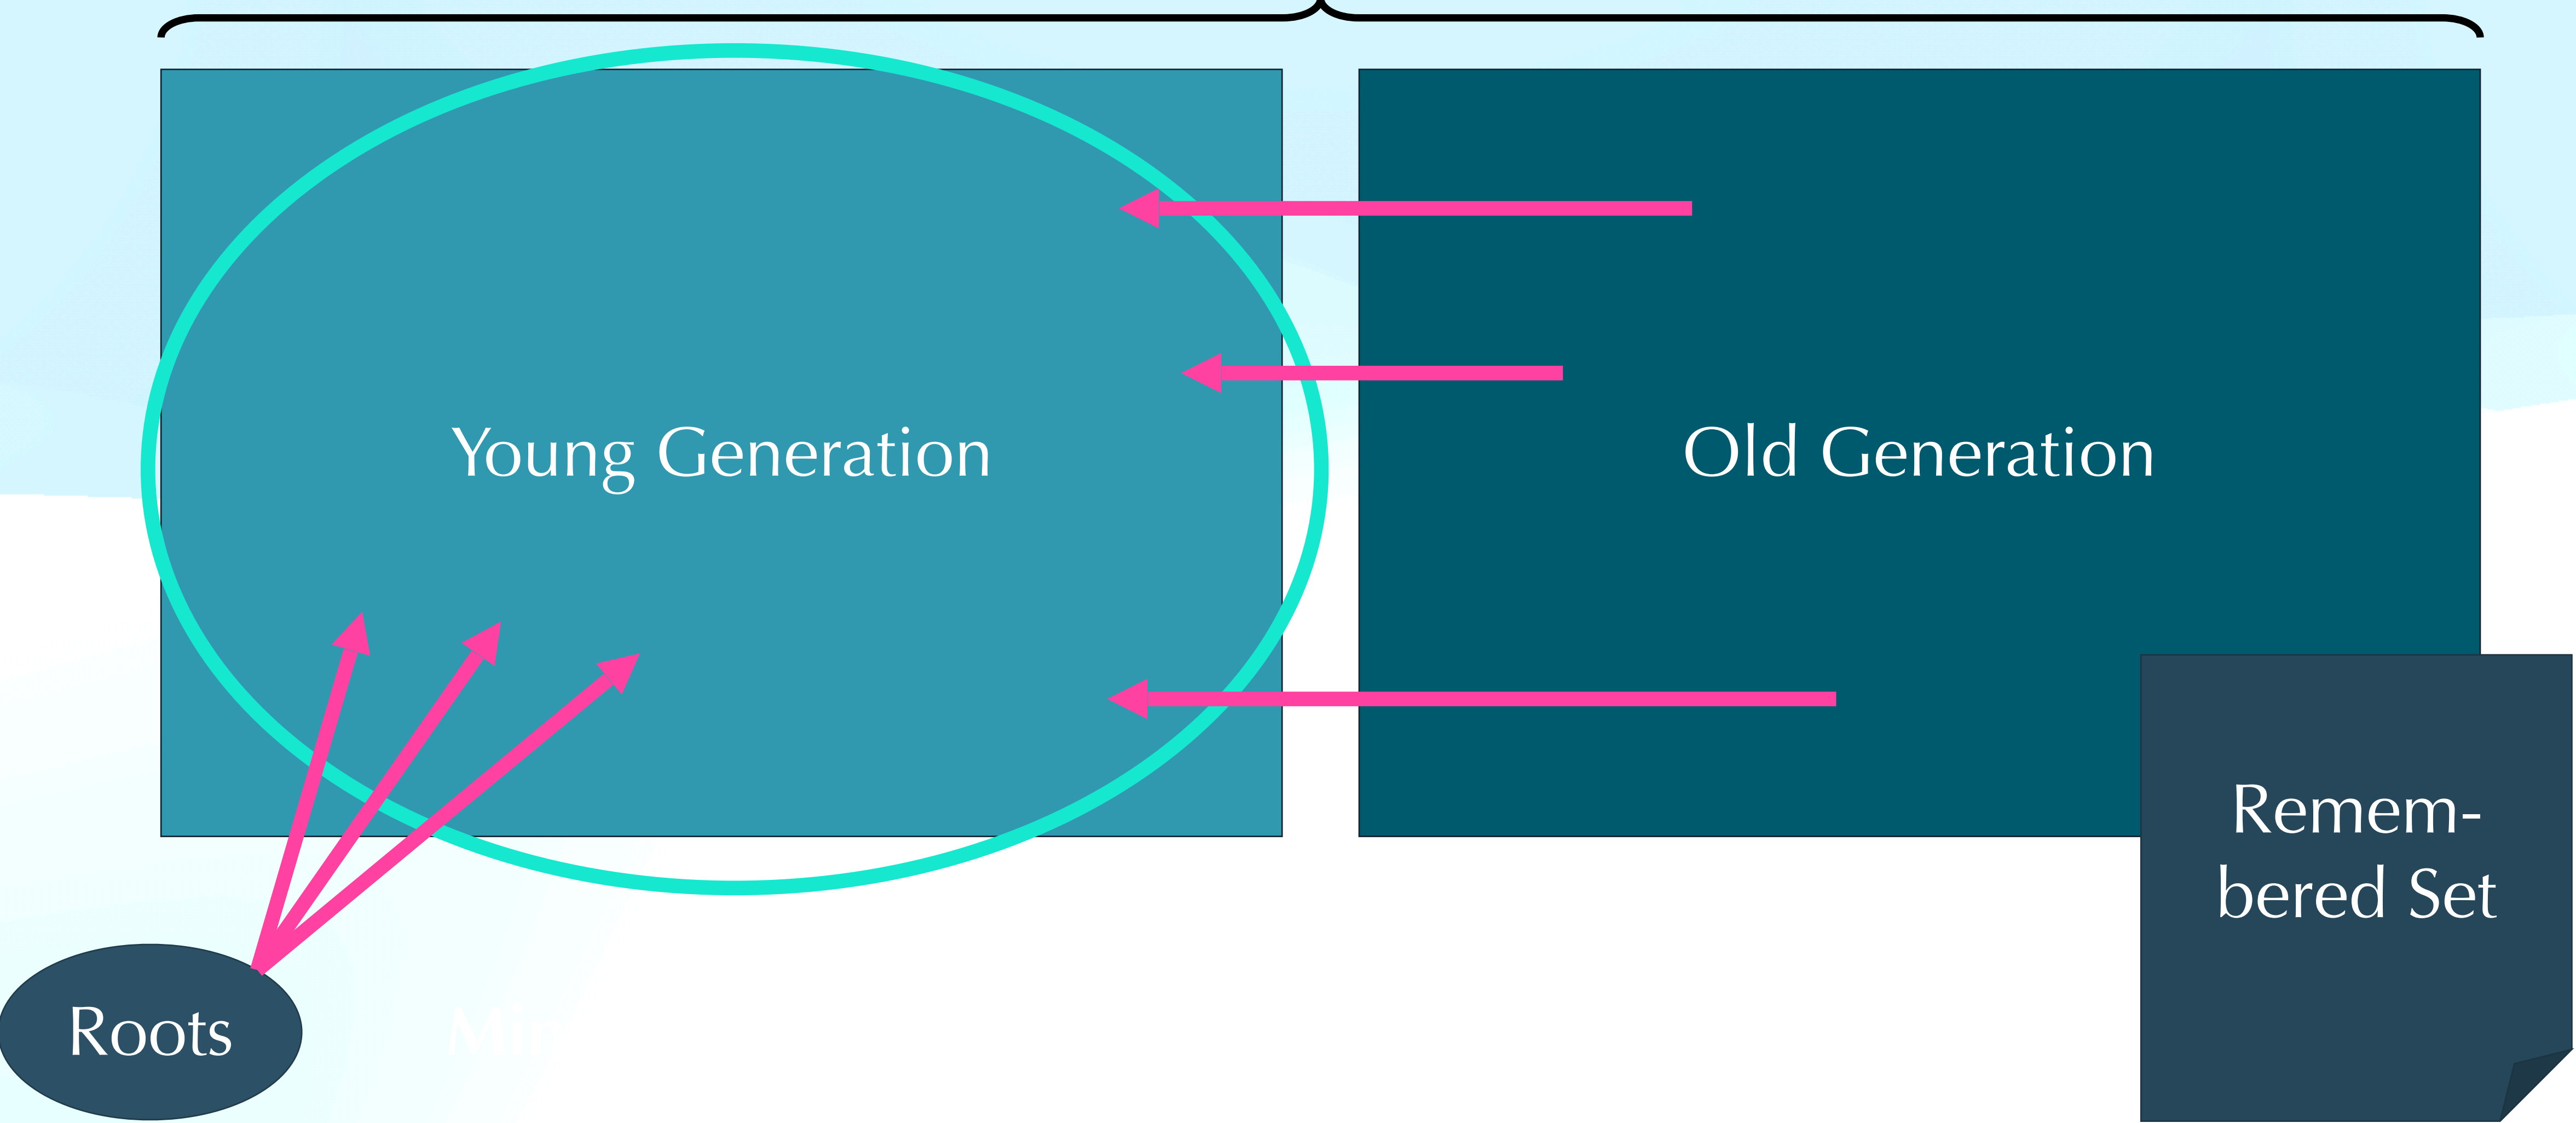


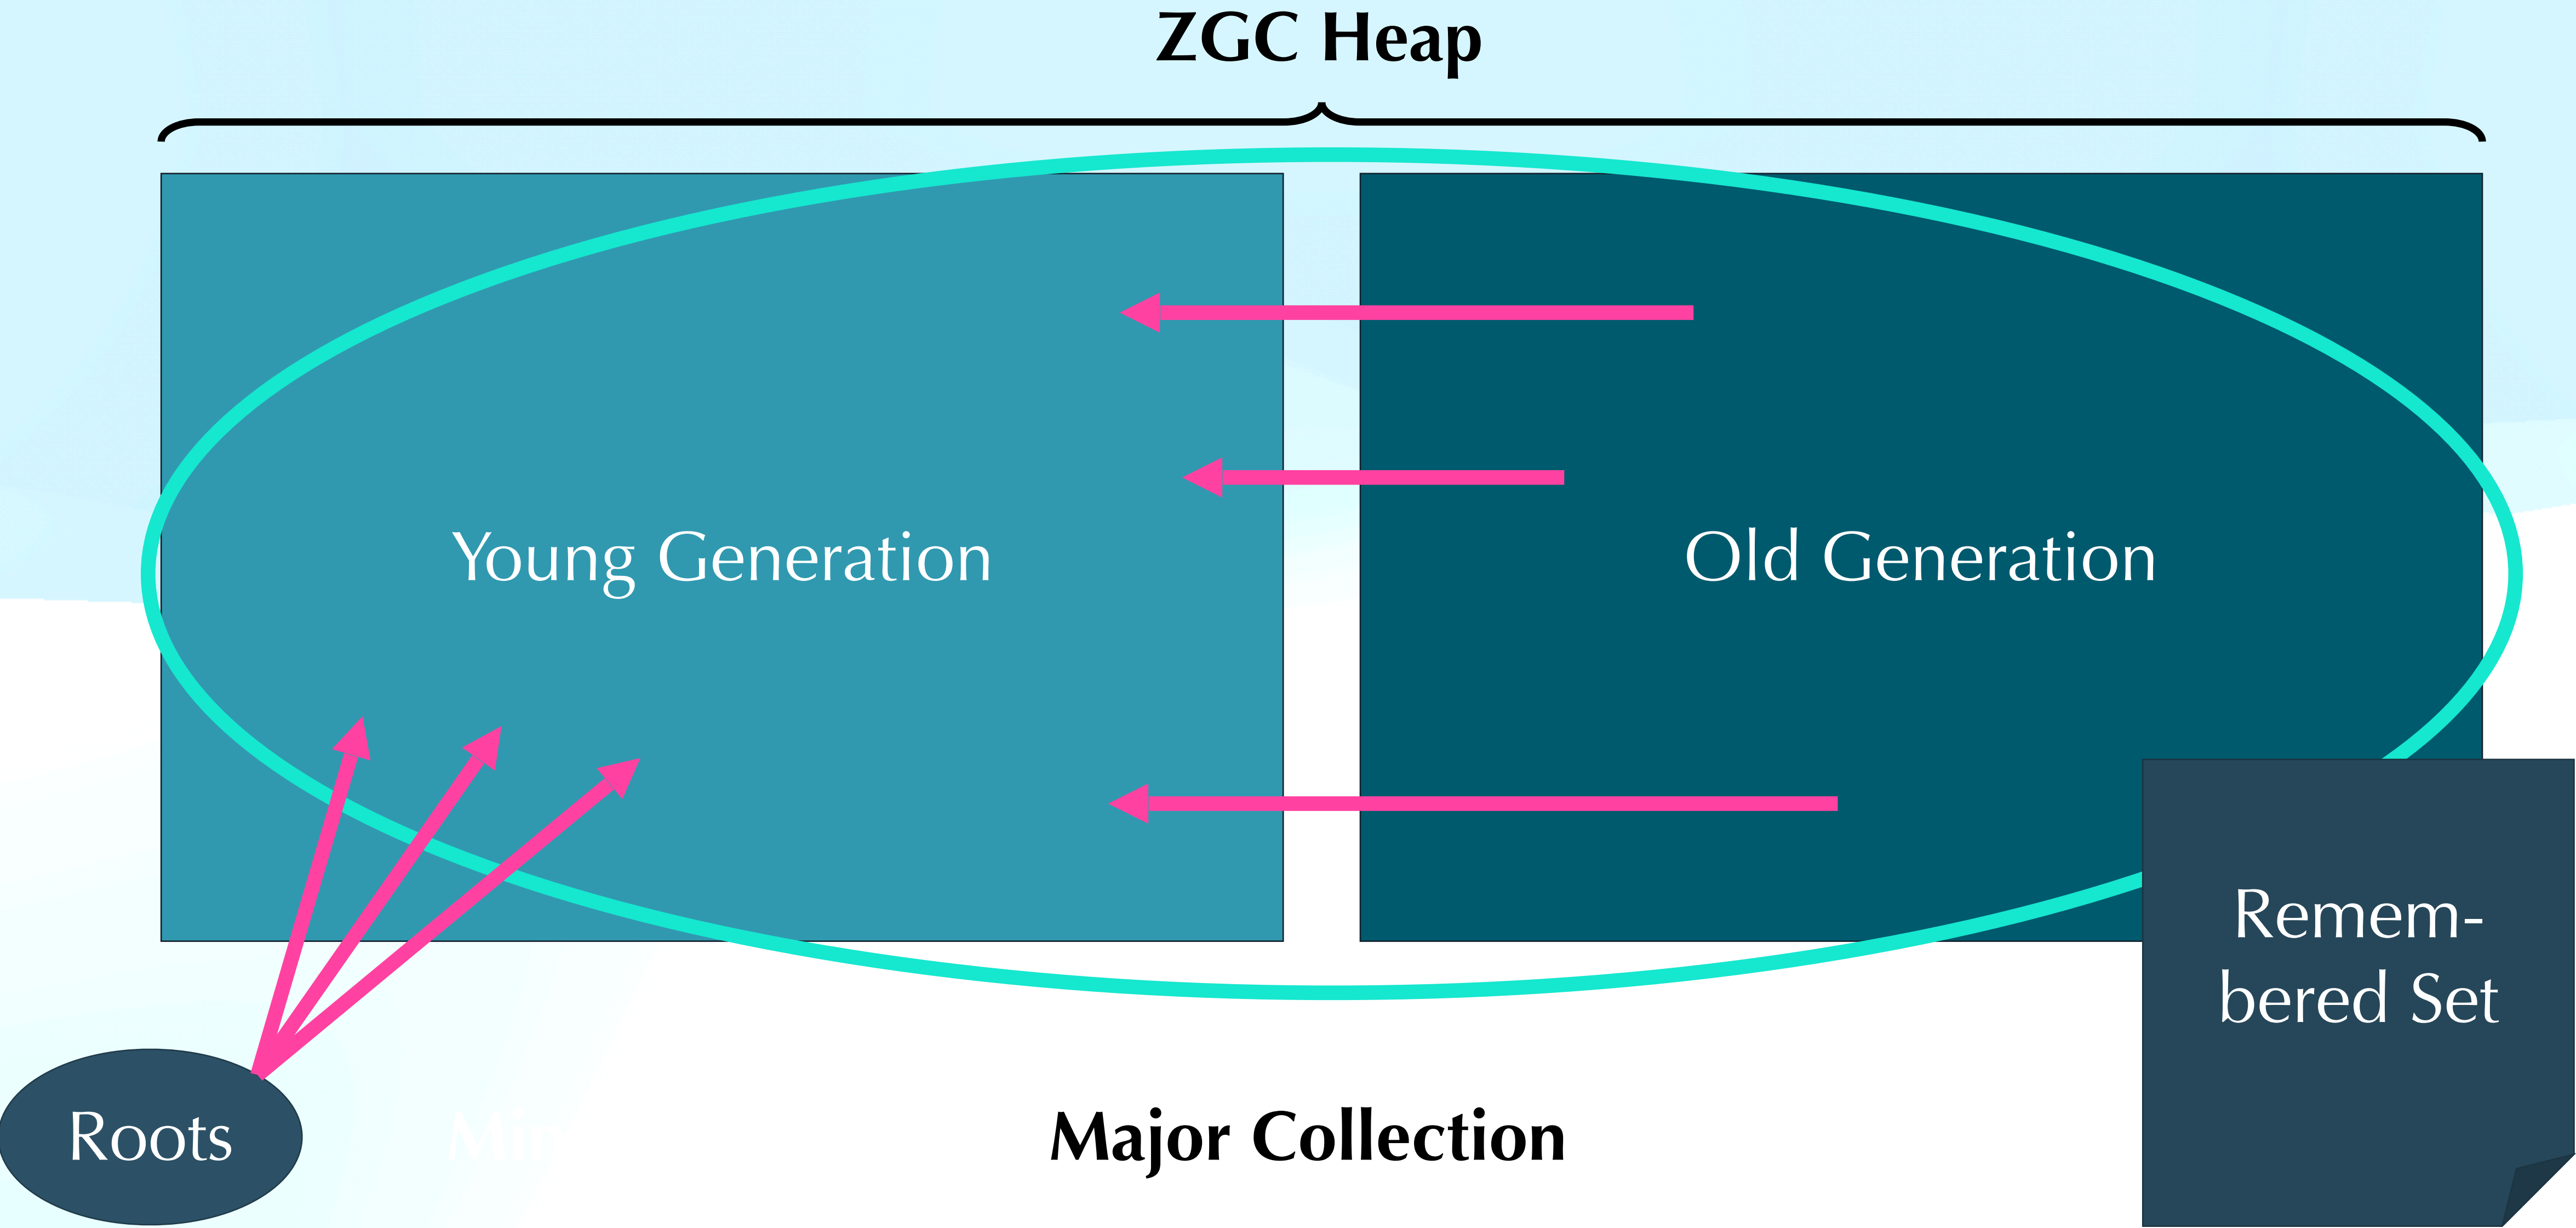
The diagram illustrates the ZGC (Zing Garbage Collection) heap structure. It consists of two main memory spaces: the Young Generation and the Old Generation. The Young Generation is represented by a light blue rectangle on the left, and the Old Generation is represented by a dark blue rectangle on the right. A horizontal curly brace spans the top of both rectangles, with the label 'ZGC Heap' centered above it. The text 'Young Generation' is centered within the light blue rectangle, and 'Old Generation' is centered within the dark blue rectangle.

Young Generation


Old Generation

ZGC Heap





ZGC Heap



Generational ZGC

- Single-generation ZGC only used load barriers
- Pointer metadata need grows (colour space); multi-mapping no longer feasible
 - But we still want to keep load barrier at two instructions
- Generational ZGC needs store barriers as well
 - Remembered-set maintenance
 - Colourless roots



Comparing Pointer Anatomies (simplified)

Single Generation ZGC



Generational ZGC



Colourless pointer



Colour Hierarchy



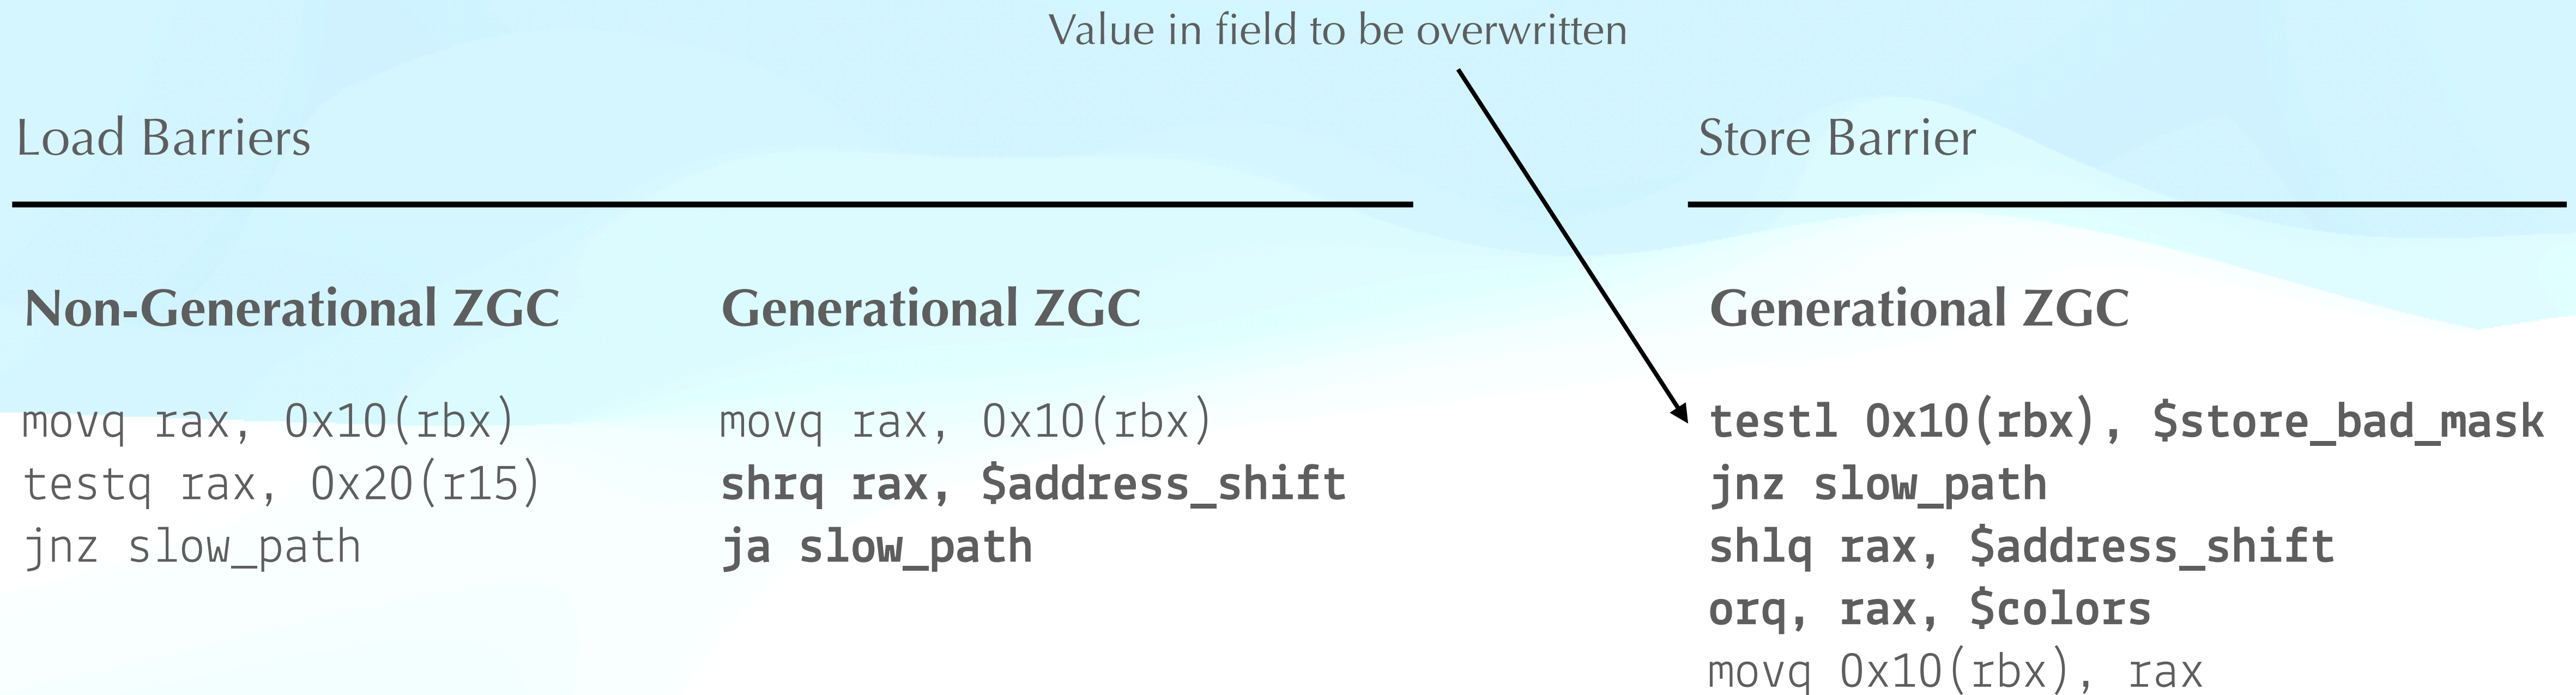
Load Good < Finalizable Good < Mark Good < Store Good

OK to load

Marked live

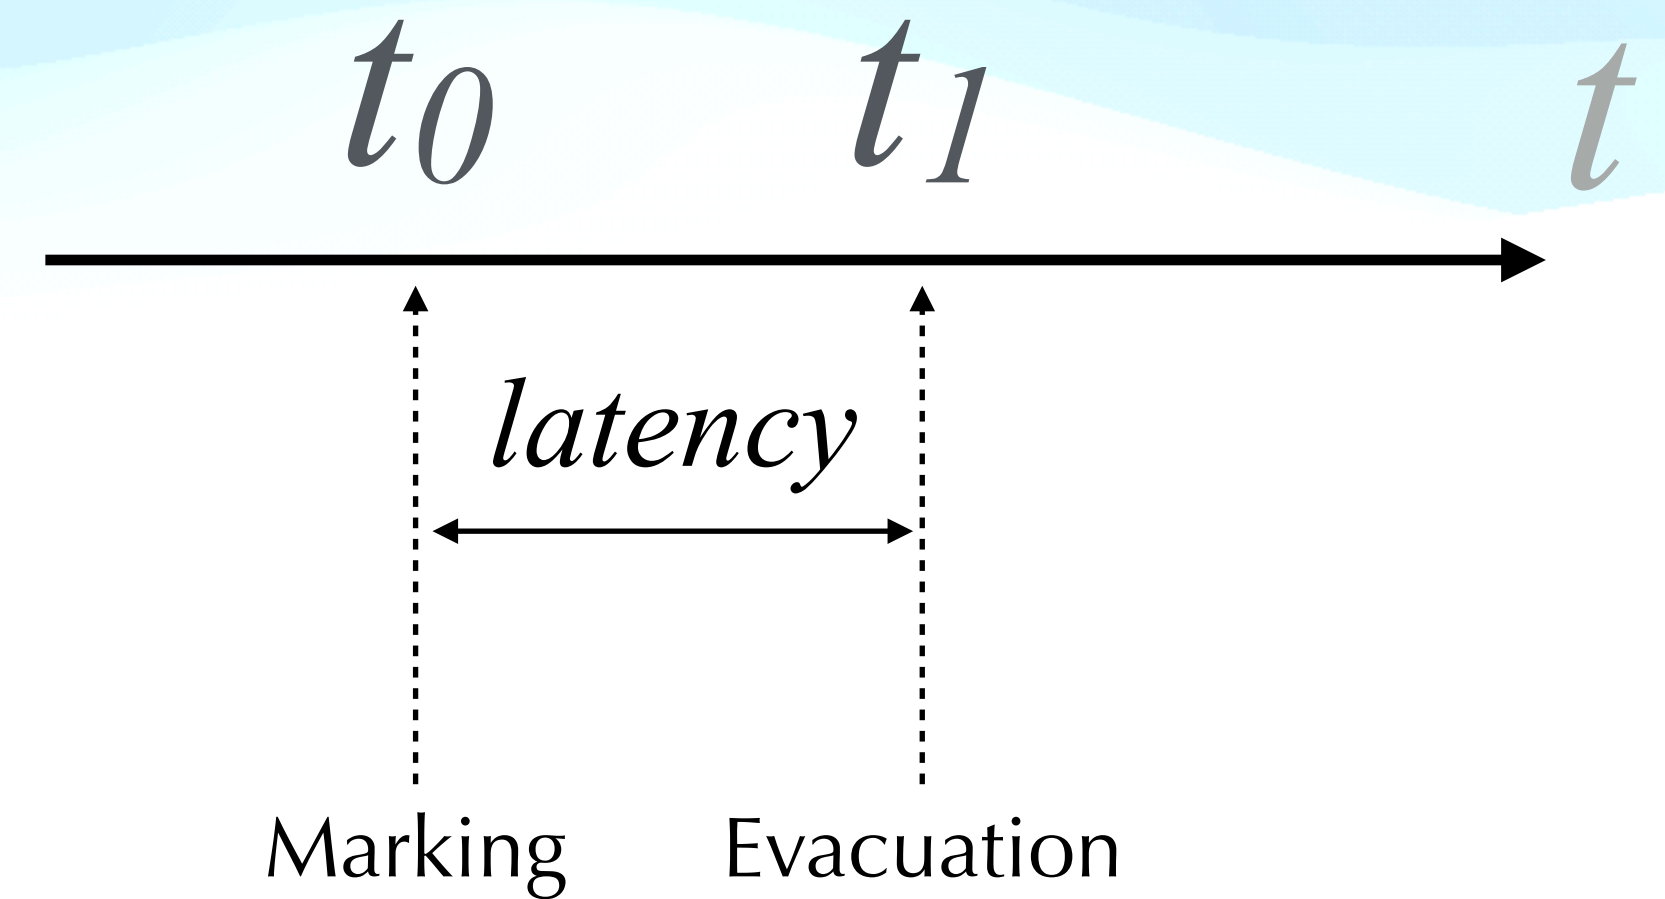

Added to remset

Evolution of ZGC Barriers



- \$address_shift, \$store_bad_mask, and \$colors are updated with cross-modifying code
- Carry Flag = last removed bit
- Zero Flag = is it null
- “Jump if above” jumps iff CF == 0 && ZF == 0


Snapshot-at-the-beginning marking



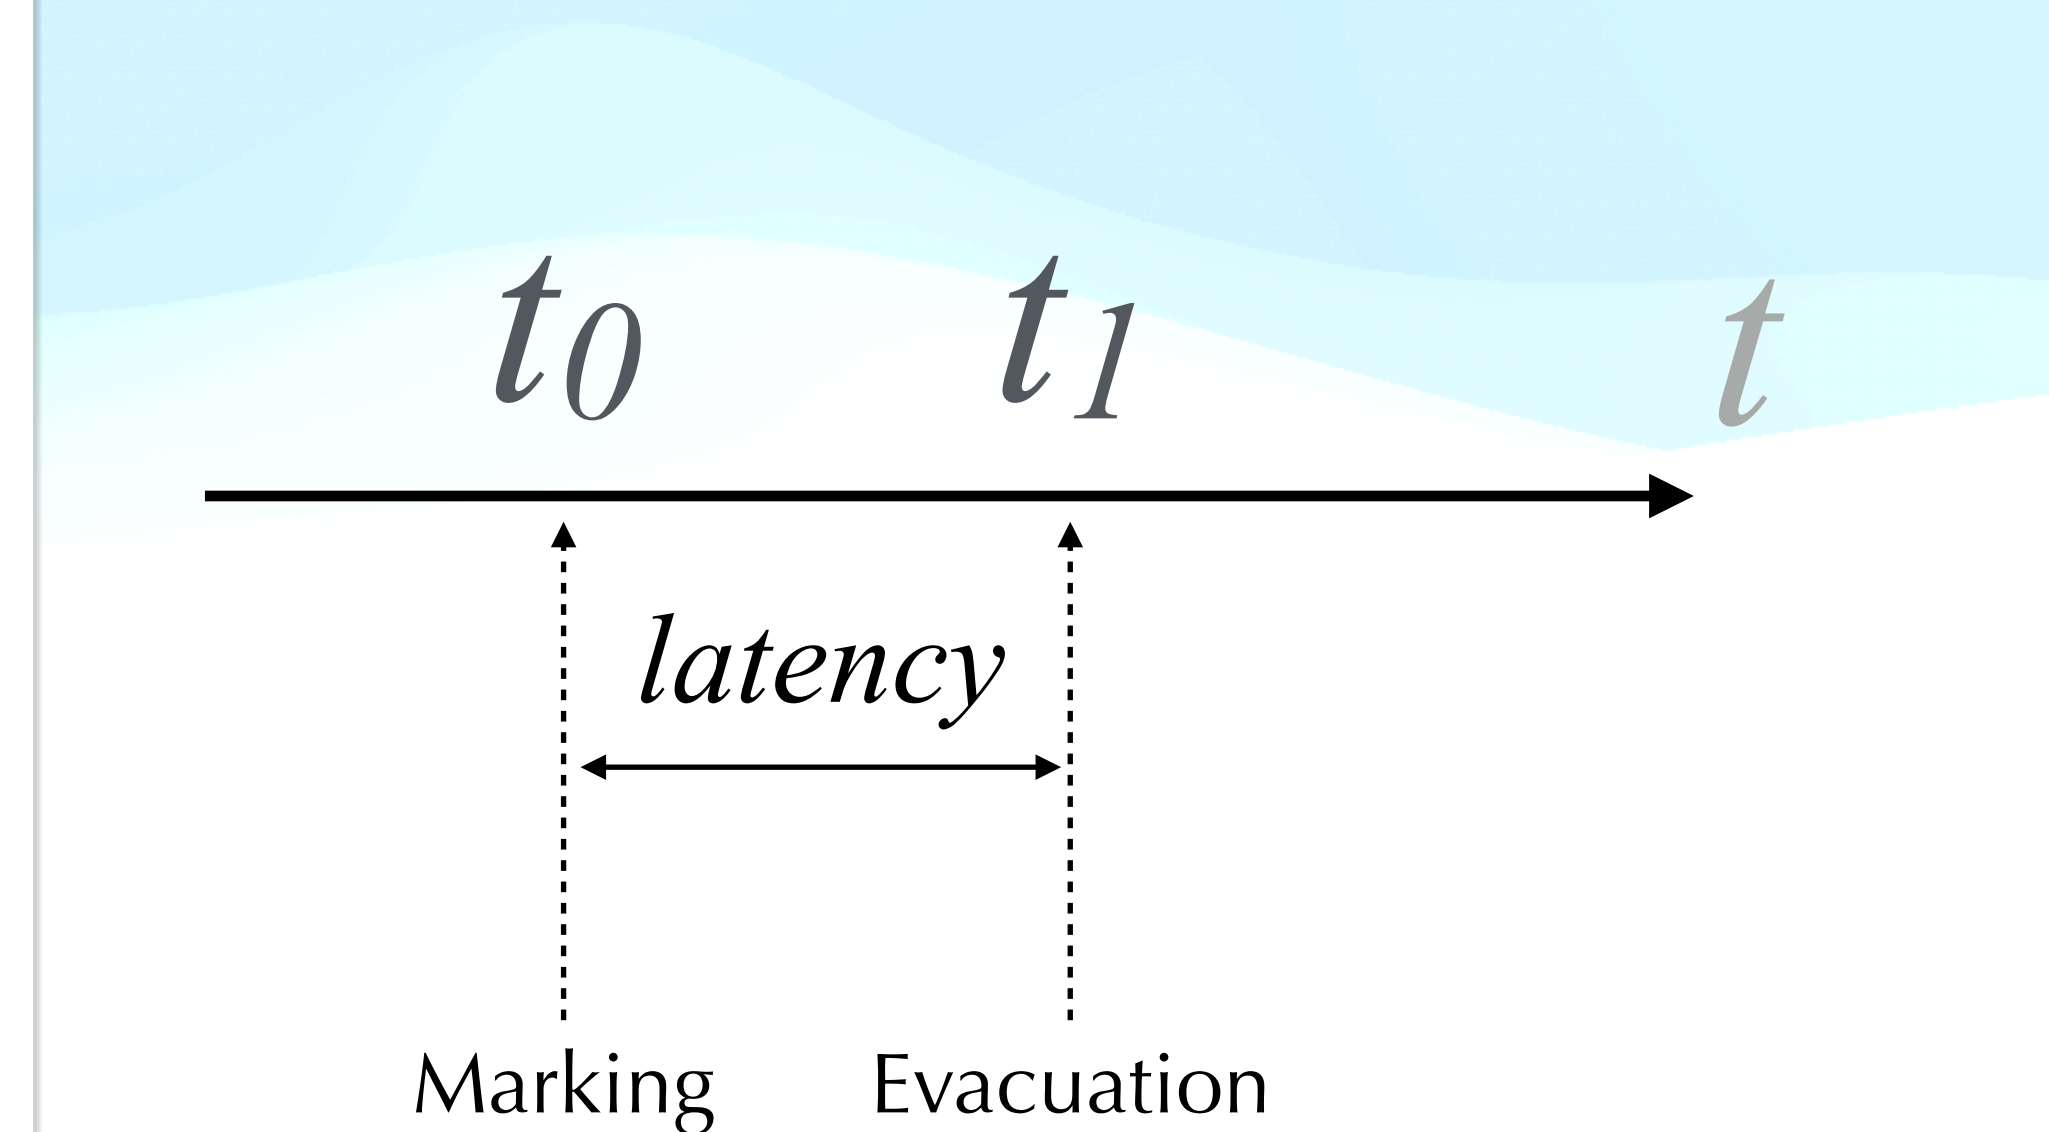
How good is information from t_0 at t_1 ?

Consequences of acting on stale liveness

Investigation using DaCapo (version at Git hash 6e411f33)




25% of all objects that we copy were effectively dead at the time



Consequences of acting on stale liveness

Investigation using DaCapo (version at Git hash 6e411f33)



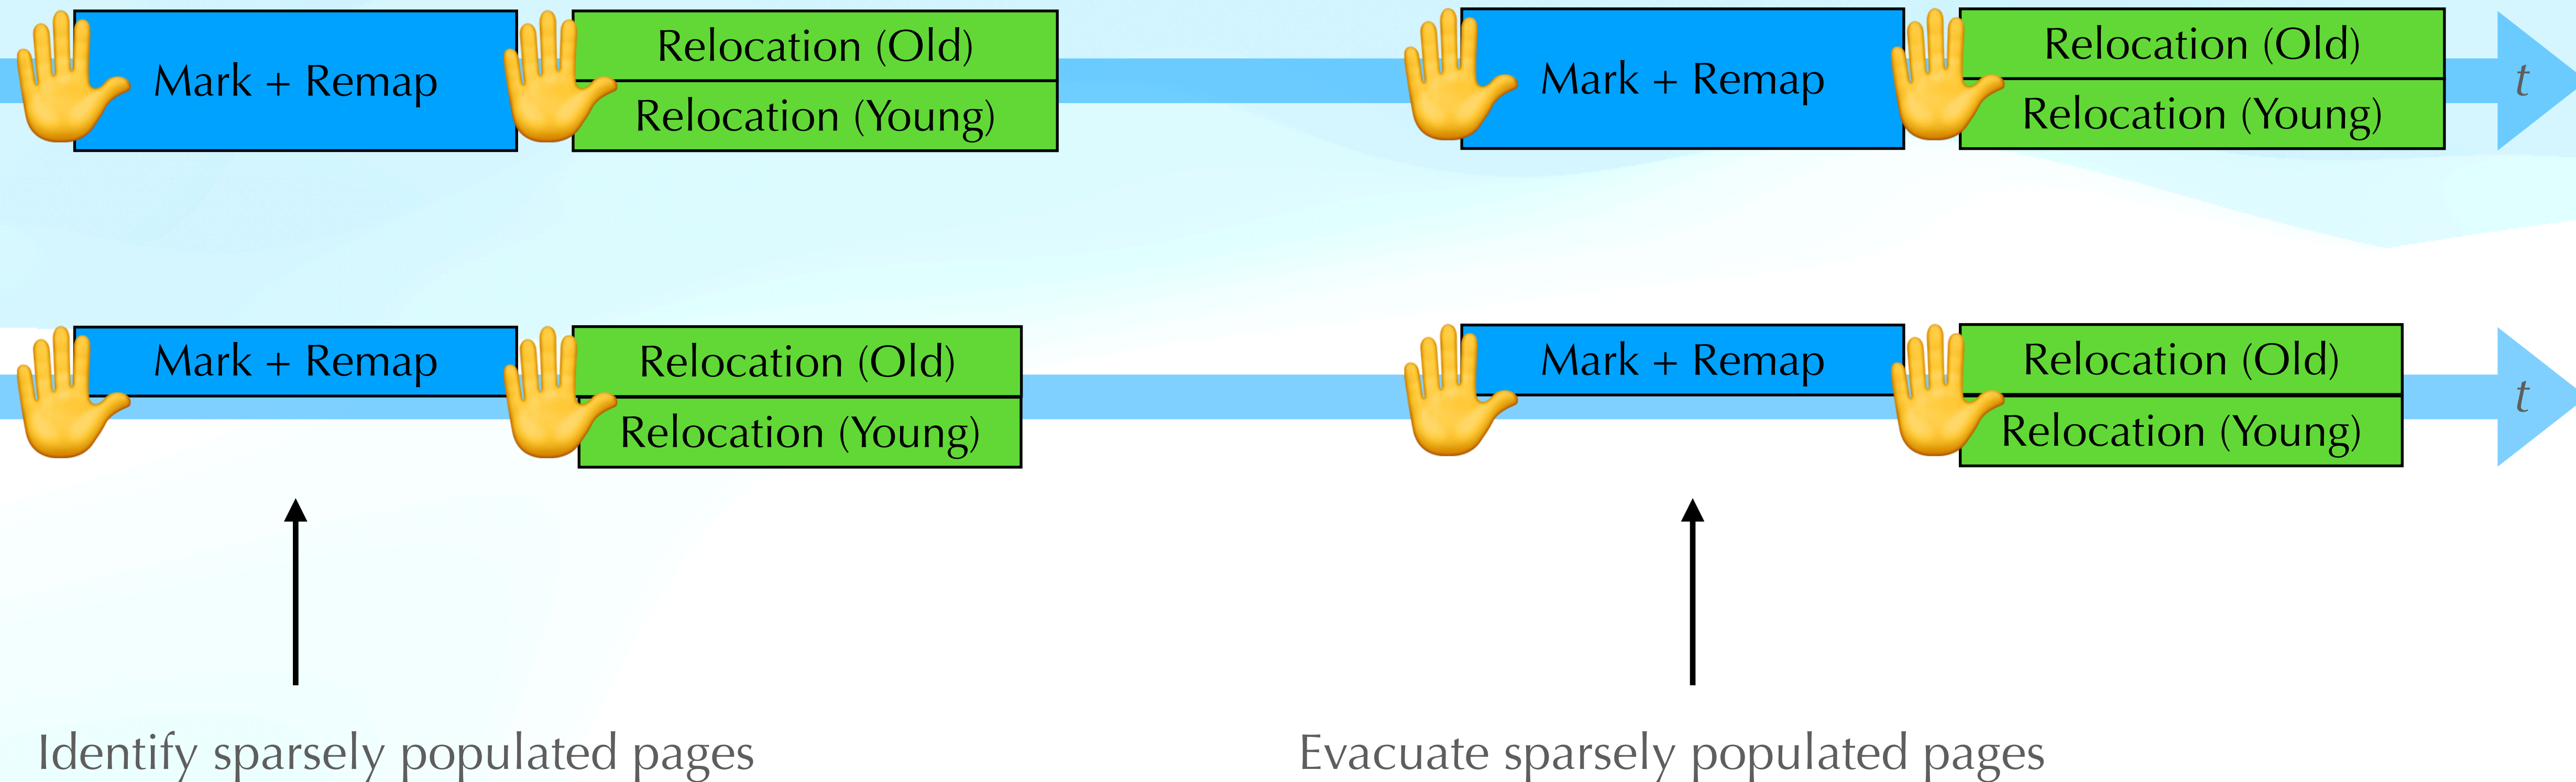
Spring	GB	%
Relocated	398	100 %
Accessed	31	8 %
Dead	367	92 %



25% of all objects that we copy were effectively dead at the time

Copy fewer dead objects & reduce barrier storms


Disclaimer: work-in-progress



Delaying relocation means more objects will be dead at relocation time

Insanely preliminary results

- Scavenging but only on sparsely populated pages
- Young relocation in some graph order
- Reduce CPU usage of ZGC
- Flatten the curve of mem. reclamation
- Challenging to work with the 12 colour bits
- Unclear what the good heuristics are



10 iterations of Spring — $\approx 3.6\%$ speedup

Some Concluding Remarks

- Load barrier-based designs are everywhere and they can work well
- Generational ZGC complexity is higher than non-generational
- Load barriers are very useful to have in your language
 - Concurrent relocation
 - Locality optimisations
 - Various forms of telemetry
- Would be very nice to close the performance gap between concurrent collectors and throughput collectors because the design is (conceptually) simpler